

# Configuring autonomous vehicles

Dr. Silviu S. Craciunas  
Principal Scientist

## Copyright Notice

All Rights Reserved.



Unless otherwise stated, the property right for the document is owned exclusively by TTTech Auto AG. Regarding all rights (e. g. possession, copyrights), no license or any other right is given to anyone. You may not copy, reproduce, distribute, publish, display, perform, modify, create derivative works, transmit, or in any way exploit any such content, nor may you distribute any part of this content over any network, including a local area network, sell or offer it for sale, or use such content to construct any kind of database. You may not alter or remove any copyright or other notice from copies of the content. Copying or storing any content except as provided above is expressly prohibited without prior written permission of TTTech Auto AG or the copyright holder identified in the individual content's copyright notice. For permission to use the content, please contact **[silviu.craciunas@tttech.com](mailto:silviu.craciunas@tttech.com)**.

# Project highlights



TTTech Auto safety and robustness in series production

Hardware development until the A-Sample grade



Scalable distributed control system from TTTech Industrial into Vestas 2, 4 and 6 Megawatt turbines



Collins Aerospace Power System and Air Management for the Boeing 787-8, 787-9, 787-10 family



TTTech TTEthernet Switch and End System IP products for the avionics system.

TTEthernet is the “nervous system” i.e. avionics network platform of the space craft



Accelerating  
software-defined  
vehicles together





# Introduction



# Autonomous vehicles

Increase in Safety – Fewer Accidents


Reduced Emission (The use of electric autonomous taxis alone could reduce greenhouse gas emissions by 90% per mile by the year 2030)

– Lower Fuel Consumptions

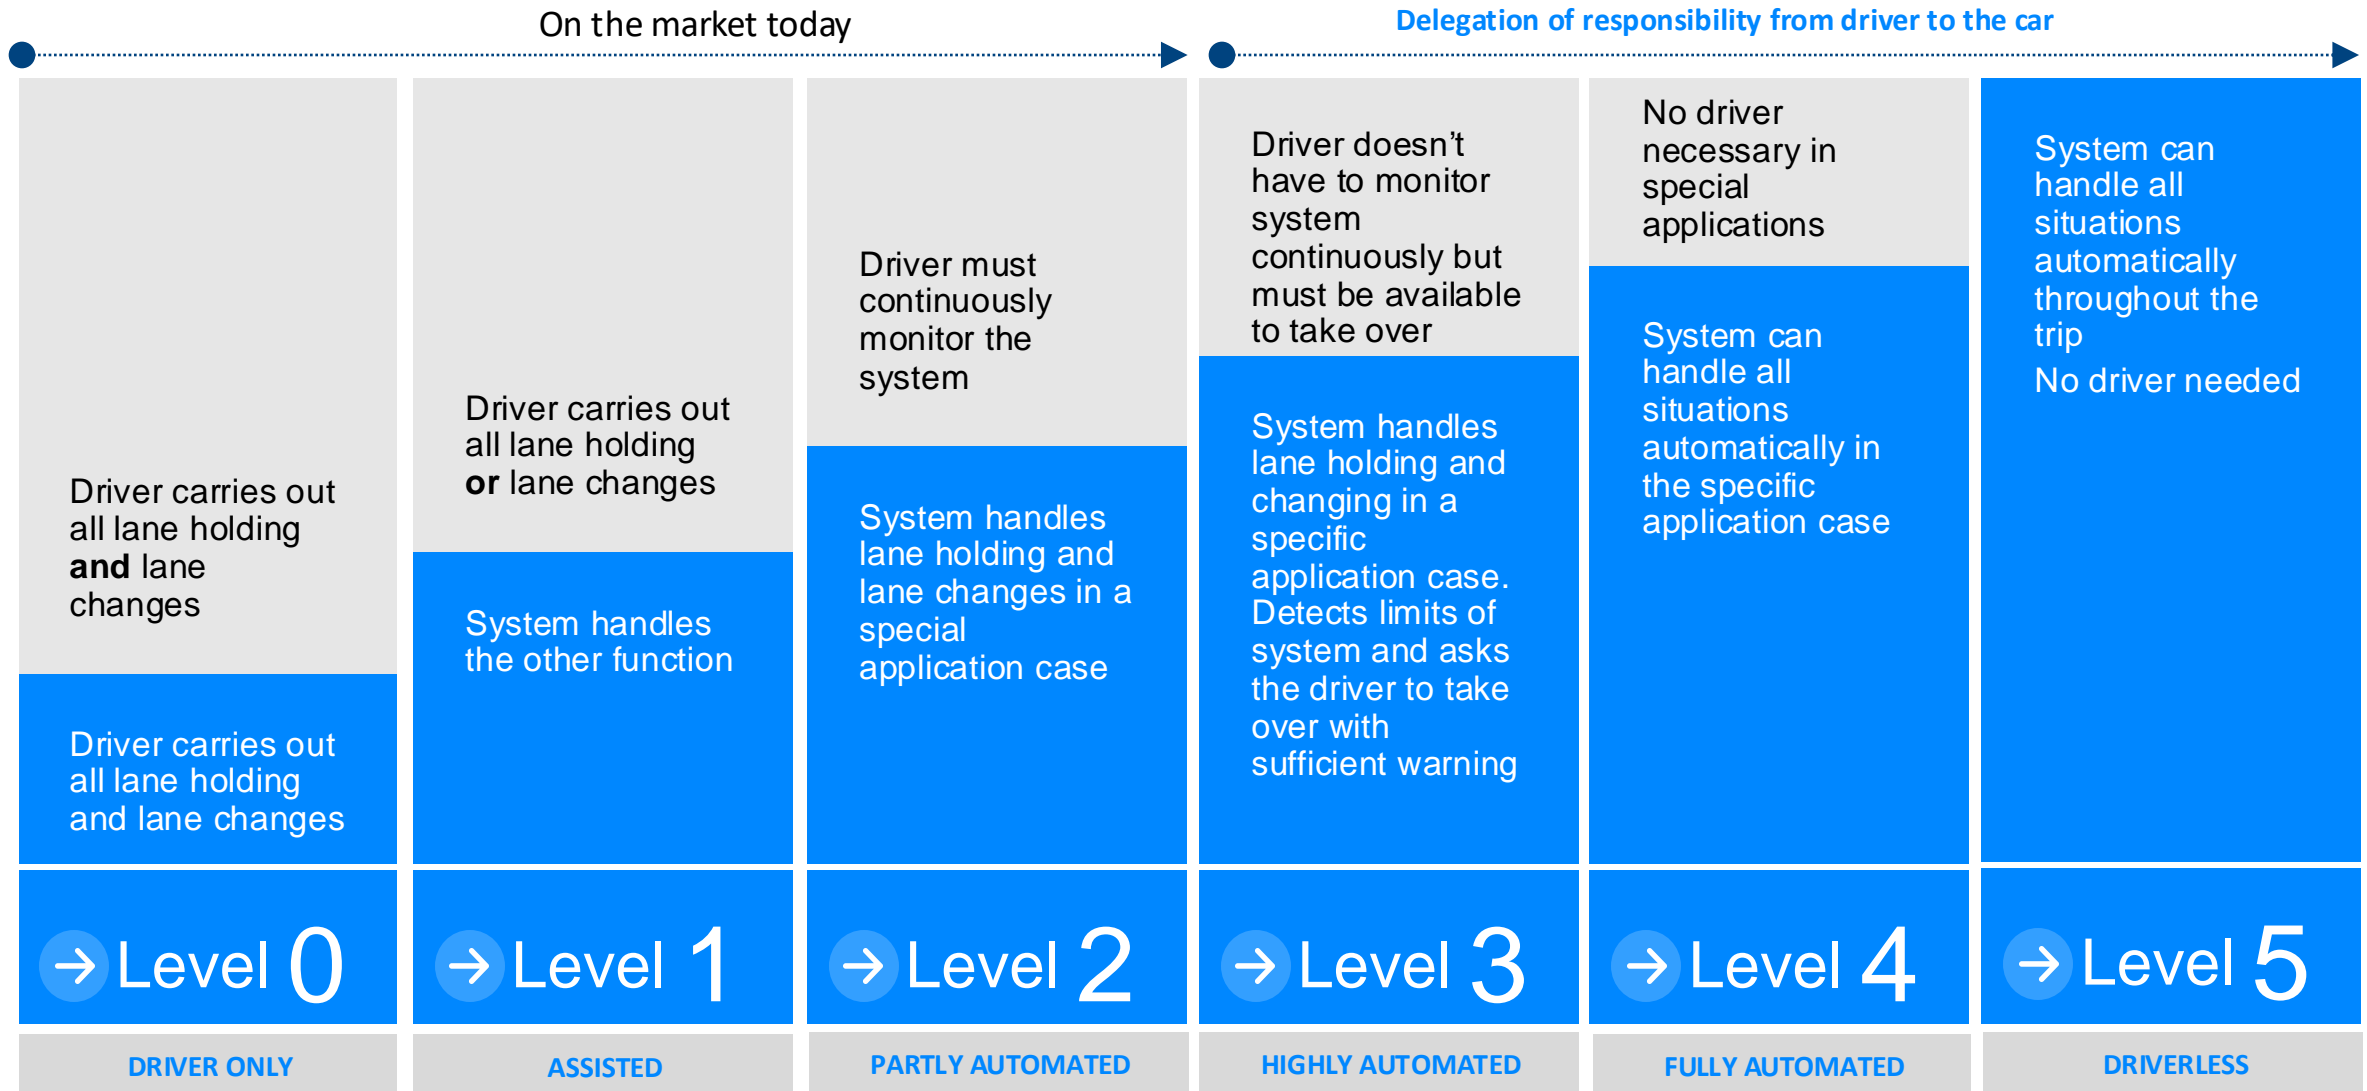
Better Health – Less Stress Through Traffic Jams

– Much Better Predictability

Less Traffic Jams – Increased Highway Capacity



# Automated Driving - System Classification by VDA



# Challenge: The Fail-Operational Requirement

Level 3


“System detects its limits and asks driver to take over in case of faults”

Level 4

“System can handle all situations automatically in the specific application case”

**Already from Level 3 onwards there is a requirement for fail-operational behavior. Component failures must be tolerated!**

- Continue operation and request driver to take over (~10s in Level 3 systems)
- If driver does not take over, then situation specific reaction e.g. controlled stop



Change lanes and stop the car at safe position



Stop the car in lane (avoiding obstacles)

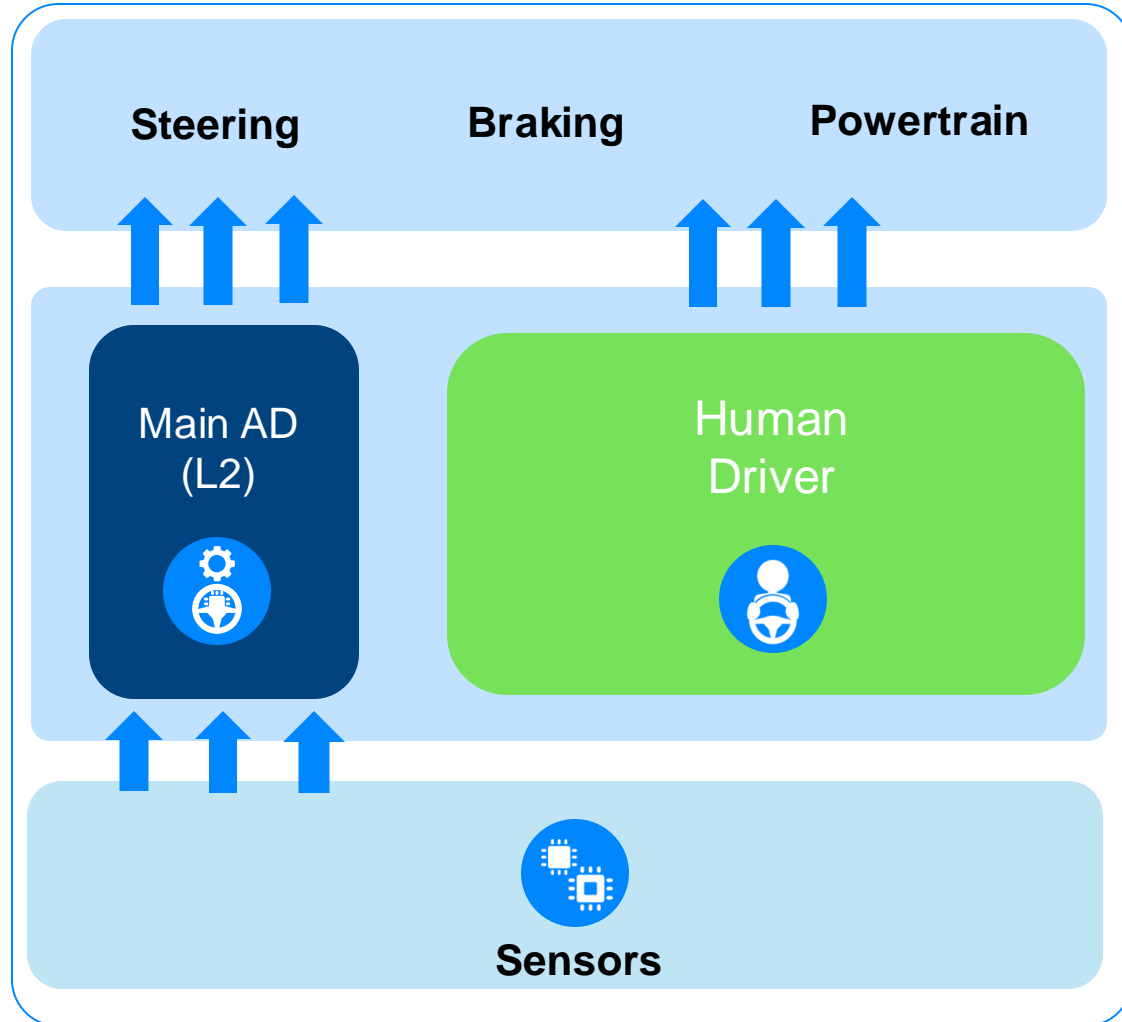
Level 5

System has to continue driving until planned travel end point

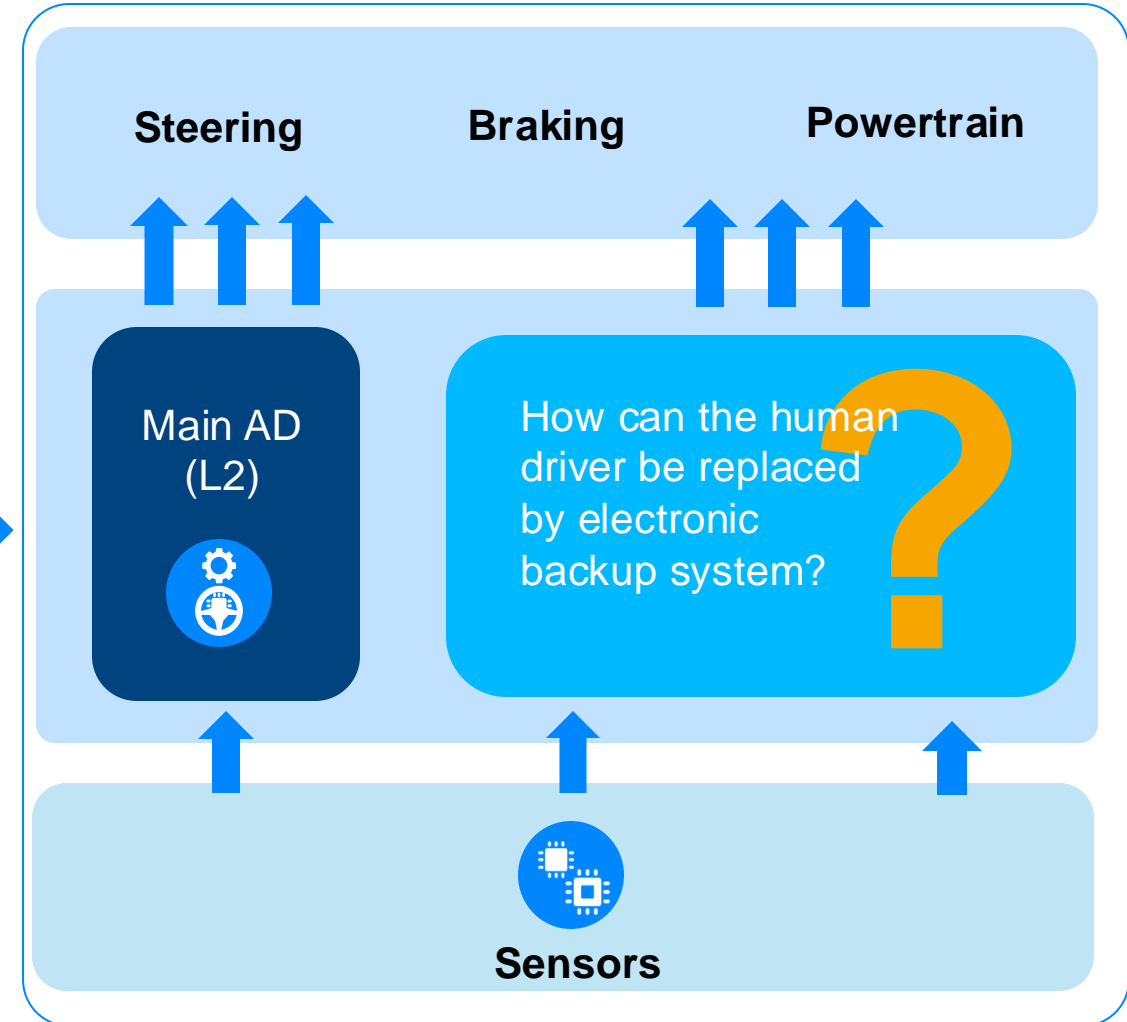


# Moving from a Level 2 to a Level 3/4 system

Level 2 system (fail-silent)



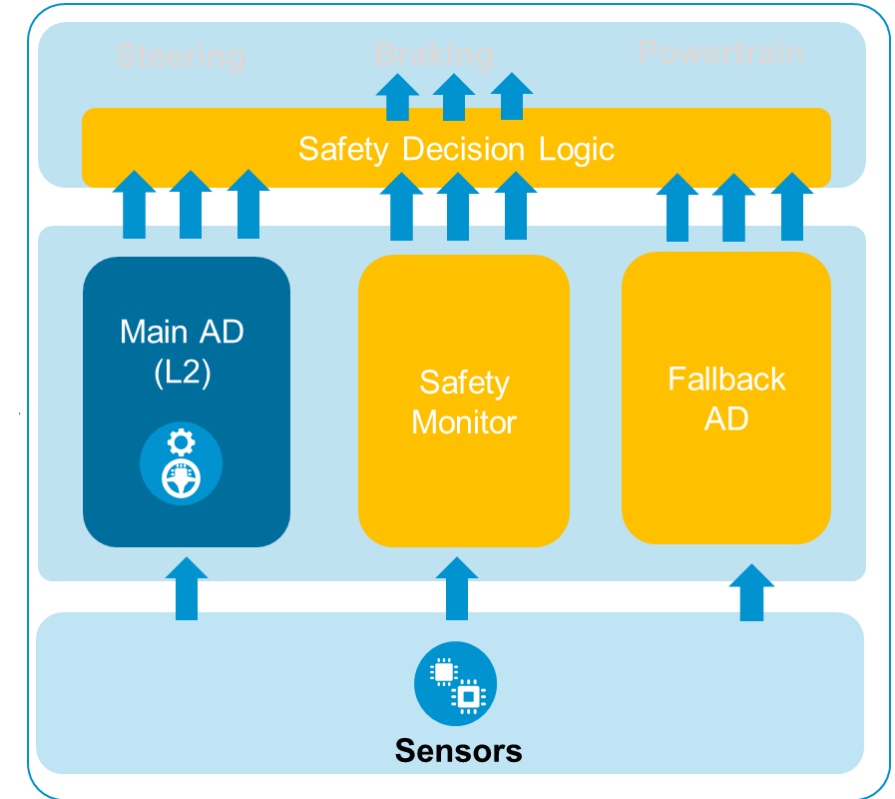
Level 4 system (fail-operational)



# Conceptual Architecture

## Systematic partitioning to solve key „impossibilities“:

1. Impossible to **find all design faults** in a large and **complex monolithic software** system
2. Impossible to **avoid single event upsets** (e.g., bit flip) in **non-redundant hardware** during the life-time of an ultra-dependable system
3. Impossible to **establish the ultra-high dependability** of a large monolithic system **by testing and simulation**
4. Impossible to precisely **specify all edge cases** that can be encountered in driving situations
5. Impossible to source AD system modules with full ASIL D qualification: SOCs, OSs, BSW, ... contain lower ASIL and QM elements



In order to realize automated driving at higher levels tens to hundreds of functions (applications) need to be integrated with each other to operate as a coordinated whole.


# Modern ADAS/AD paradigm

Growing number of complex ADAS functions

- automated/assisted parking, lane changing, emergency brake assistance, autonomous driving.


Fusion of multiple software functions into the same hardware platform

- reusability and portability
- integration of mixed-criticality functions with guaranteed temporal and spatial isolation




**n ECU: n functions**

DISTRIBUTED  
E/E ARCHITECTURE



**5-7 domains: n functions**

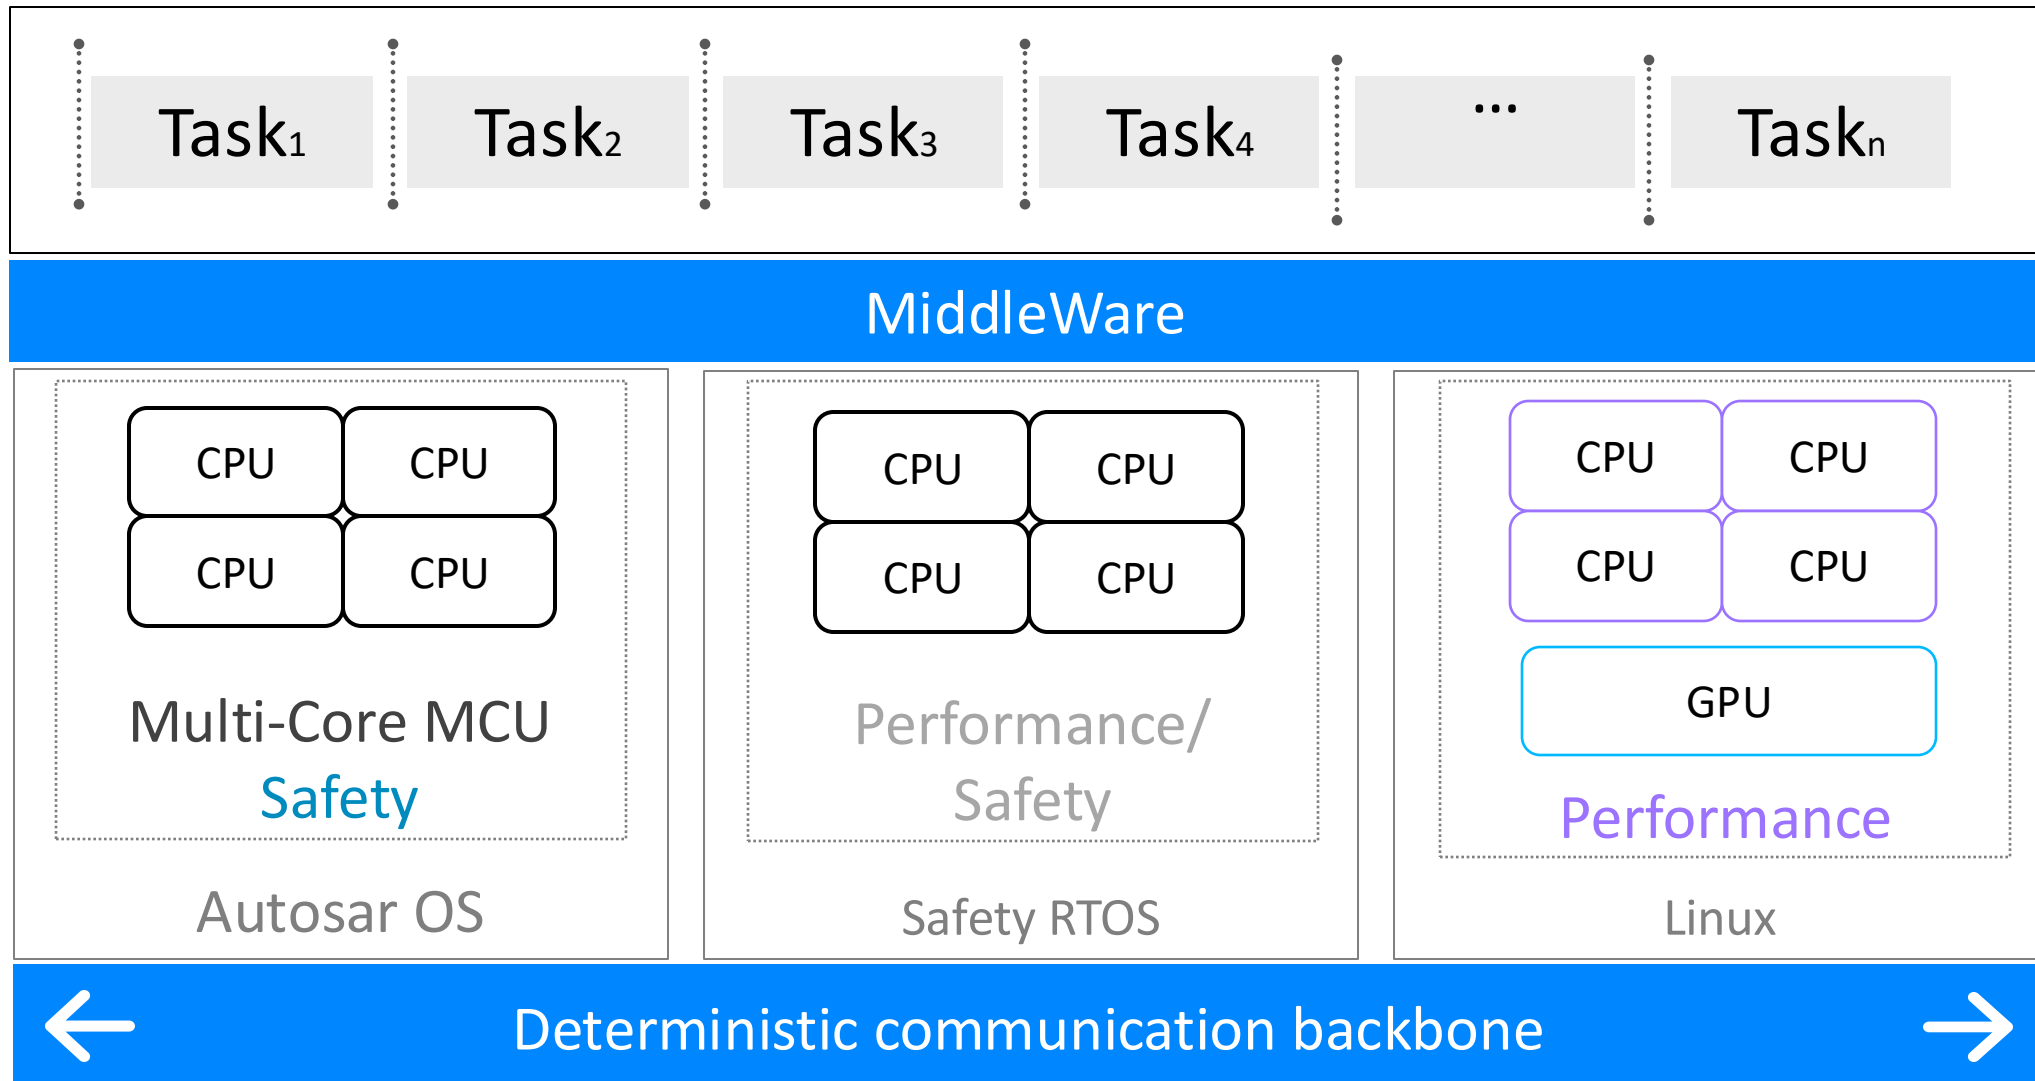
DOMAIN  
E/E ARCHITECTURE



**2 ECUs: n functions**

CENTRALISED  
E/E ARCHITECTURE

# ADAS/AD platform



# ADAS/AD systems are real-time systems

AD systems (and real-time systems in general) are reactive systems.



Given an event, the system must produce a reaction or response within a specified time (deadline).

Periodic events occur at predictable (regular) time intervals

- deterministic – all relevant information is available at design time
- complex requirements - deadlines, end-to-end latency, jitter, ...

Sporadic events occur with a certain maximum rate

- usually defined by minimum interarrival time

Aperiodic events occur at unpredictable times

- variance – not everything is completely known, especially arrival times
- limited predictability: only simple requirements can be guaranteed in the best case

Most real systems  
contain periodic,  
sporadic and aperiodic  
events



# Sensing, computing, actuating



Sensor fusion + scene segmentation based on deep neural networks

Environment camera



Laser scanner



Stereo camera



Ultrasound sensors



Lidar



Radar



Front camera



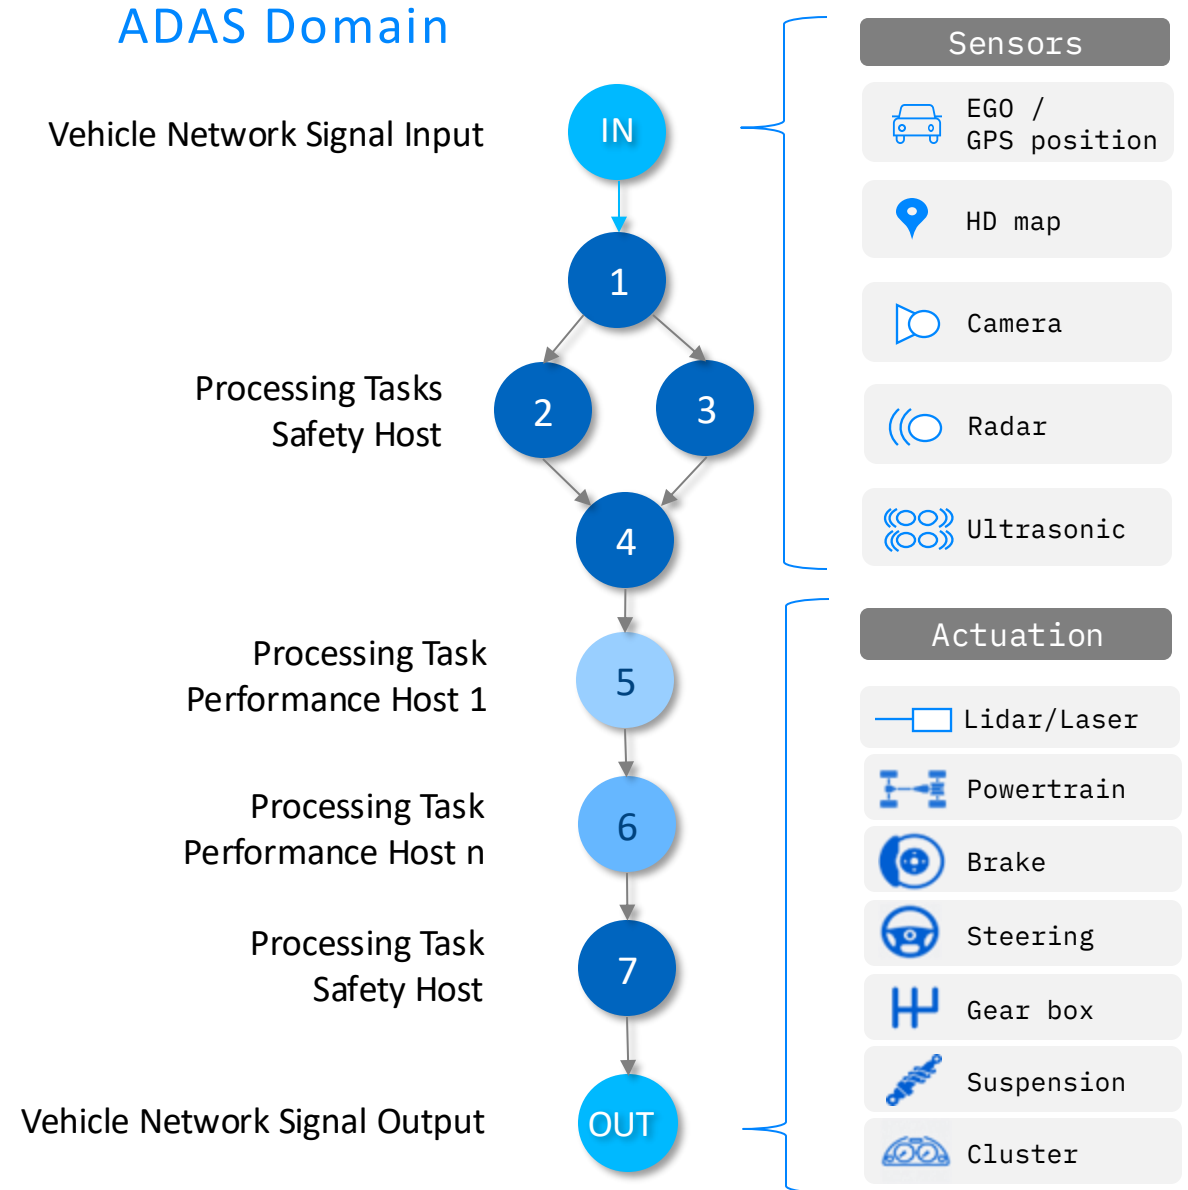
# Cause-effect chains

Cause-effect chains are a characteristic of automotive software:

- provide additional timing and dependency requirements on the execution of tasks
- can span across multiple activation patterns
- include multiple tasks, even the same task multiple times
- have priorities and end-to-end latencies
- include communication latencies

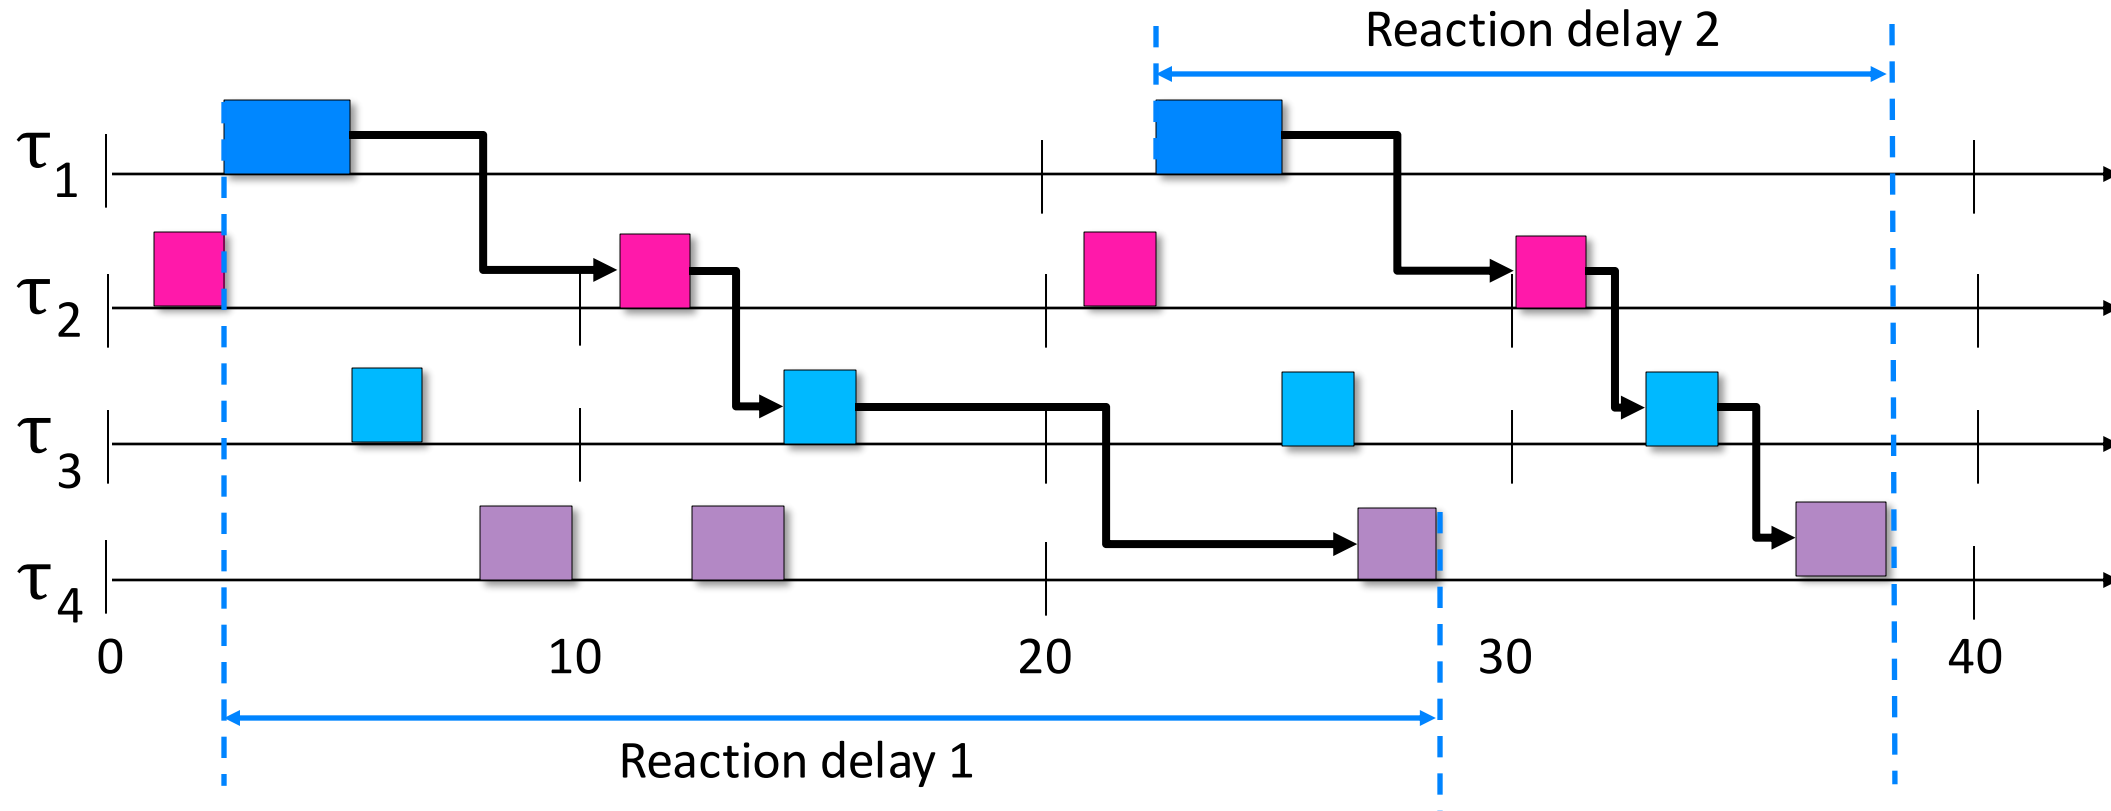
Semantics:

- Reaction delay
- Data age
- Order



# Cause-effect chains: reaction delay

After each event, there must be a reaction within the maximum allowed time





# Correct-by-design approach

# Correct-by-design approach



## Correct-by-testing

- Manual resource allocation and priority
- Extensive testing
- Empirically driven
- No guarantees for safe and reliable runtime execution in corner cases

No longer address AD central  
compute complexity



## Correct-by-design

- ✓ Automates manual process
- ✓ Hardware resources are modeled
- ✓ Mathematically proven approach
- ✓ Addresses all corner cases
- ✓ Based on very advanced heuristics
- ✓ Scales with complexity

Future-proof and  
safe solution

*“Design approaches not based on mathematically proven real time scheduling properties are prone to missing deadlines during unusual operational conditions”*

UL4600 Standard for Safety for the Evaluation of Autonomous Products



# Scheduling (dispatching) algorithms

FCFS

(Weighted) Fair queueing

Priority scheduling

(Weighted) Round-robin

...

## General-purpose computing

Main objectives: avoid starvation, responsiveness, fairness, QoS

Rate monotonic

Deadline monotonic

Earliest deadline first

Time-triggered

...

## Real-time/multimedia computing

Main objectives: meeting deadlines & other real-time constraints

# Fixed-priority scheduling

- The static priorities of tasks that are in the ready queue determine the execution order of tasks
- Priorities are usually derived from temporal requirements (mostly period and deadlines)


FP

$\tau_1(1,4)$

$\tau_2(2,8)$

$\tau_3(4,12)$

$\tau_4(1,12)$




# Utilization-based test

[Liu1973] A set of  $n$  independent, preemptable periodic tasks with relative deadlines equal to their respective periods can be scheduled under the RM algorithm if its total utilization  $U$  is at most

$$\sum_{i=0}^n \frac{C_i}{T_i} \leq n(2^{\frac{1}{n}} - 1) \quad n \rightarrow \infty \quad \lim_{n \rightarrow \infty} n(2^{\frac{1}{n}} - 1) = \ln(2) \approx 0.69$$

In practice the bound is on average 88%. [Davis2016]

When the periods are harmonic, the bound is 100%. [Mok1991]




# Fixed-priority scheduling wrt. dependency chains

High-priority task:  $T_H(1,4)$

Chain:  $T_1(2,8)$ ,  $T_2(4,12)$ ,  $T_3(2,12)$

RM




# Fixed-priority scheduling wrt. dependency chains

High-priority task:  $T_H(1,4)$

Chain:  $T_1(2,8)$ ,  $T_2(4,12)$ ,  $T_3(2,12)$

RM





# Fixed-priority scheduling

Generally, Fixed-Priority Scheduling (FPS) have the following properties:

- **Static priority:** cannot change at runtime, priorities determine the execution order of tasks
- **Responsiveness:** high-priority tasks run whenever they become active
- **Reduced latency at runtime:** if a task finishes earlier than the WCET assumption then the next lower priority task runs immediately, thereby reducing its latency
- **Scalability:** Easy algorithm that scales well with the number of tasks, proven in use
- **Non-harmonic periods:** non-harmonic periods do not influence the applicability
- **Overload scenarios:** only lower priority tasks are impacted
- **Availability:** FPS are implemented in most Operating Systems

Downside of FPS:

- **Jitter:** tasks have a high run-time variance, c.f. 3.2 in [\[Locke1992\]](#)
- **Starvation:** tasks can starve if the highest priority tasks misbehaves and always stays in the ready queue
- **Dependencies:** complex task dependencies and constraints cannot be easily supported, c.f. 4.1 [\[Xu2000\]](#)
- **Utilization Bound:** in general, the utilization bound for FPS is lower than EDF and TT, c.f. 4.2 [\[Xu2000\]](#)
- **Determinism:** the run-time variance of task execution is high leading to difficult analysis, c.f. 4.3 [\[Xu2000\]](#)
- **Compositionality:** adding/modifying tasks leads to changes in lower-priority tasks, c.f. 4.4 [\[Xu2000\]](#)

For a more in-depth comparison between FPS and TT please see [\[Xu2000, Locke1992\]](#)


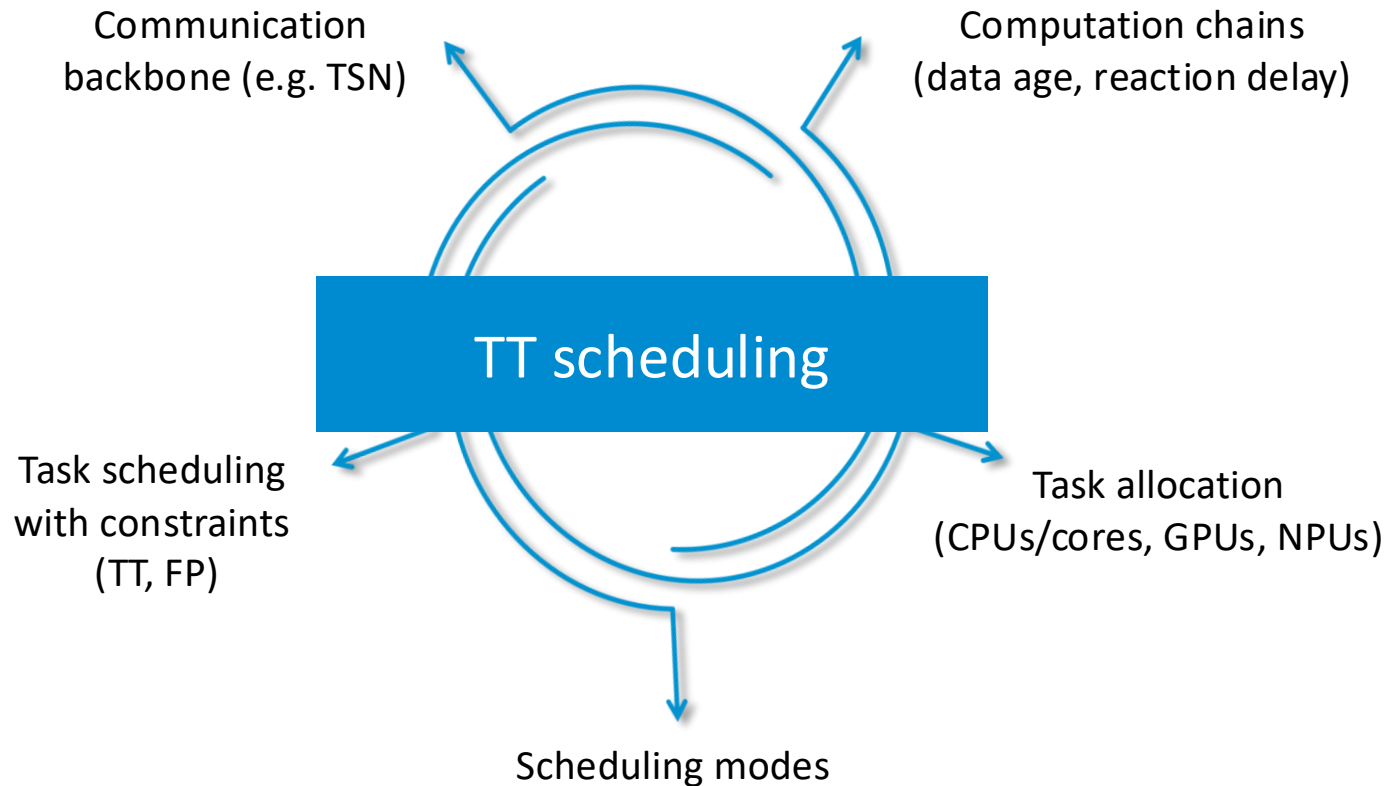
# Time-triggered vs. event-triggered

Event-based architectures have poor behavior characteristics when critical situations occur simultaneously.



# Time-triggered task scheduling

- Tasks are executed based on a **static schedule table** that has been computed at design time (offline)
- All scheduling decisions are taken at **compile time**, WCET needs to be known.
- The complete scheduling timeline for TT tasks is **fixed**.
- Simple (e.g., deadlines) and complex (e.g., chains) **constraints** are satisfied by the **schedule creation**.
- One solution is sufficient and any solution is a sufficient schedulability test.



# TT scheduling

Time-Triggered provides correct-by-design real-time guarantees:



- **Complex timing requirements:** cause-effect chains, different types of jitter
- **Temporal isolation:** no starvation possible, temporal isolation between all tasks
- **Determinism/Predictability:** **increased stability and testability**, no unwanted run-time effects, many system properties become predictable, e.g., locks, task pre-emption
- **Re-simulation:** Equivalent behaviour between cloud and embedded target
- **Synchronization to communication:** stable real-time behaviour of cause-effect chains
- **Compositionality/Integration:** adding or modifying tasks can be done incrementally, system integration of SWCs from different suppliers is easier
- **Schedulability:** more correct configurations can be realized
- **“What you see is what you get”** execution
- All critical execution and communication is deterministic and easily **traceable**.
- **Coordinated** and efficient task **execution** is enabled even beyond boundaries of single SoCs.
- Safety **certification** is simplified.

# Stability to dynamic changes at runtime

WCET is used for generating schedules and testing schedule correctness

At runtime, tasks execute for less than the WCET

Example:


BCET = 1ms

WCET = 2ms

Microtick = 1us



1000 possible execution patterns



What happens if one task executes for less than the computed WCET?




# TT vs. FP/EDF

What happens if at some point in time, one task executes for less than the computed WCET?


## Non-TT (FP/EDF)

- changes the execution of all other tasks
- state change for the entire system



## Time-Triggered

- execution of all other tasks remains the same
- no state change





## State space explosion




#tasks that run for  
less than the WCET

### Non-TT (FP/EDF)






TT



2




3



4



...




# State space explosion

How many system states do I have to potentially check/look at to give design-time guarantees?

## Non-time-triggered (FP/EDF/etc.):

- changes the execution of all other tasks
- state change for the entire system



Number of possible assignments of tasks to cores

Number of priority orderings

Number of system states

$52^{400}$

$400!$

$10^{1200}$


$= 10^{2755}$

# State space explosion

How many system states do I have to potentially test/look at?

## Non-time-triggered (FP/EDF/etc.):

- changes the execution of all other tasks
- state change for the entire system



## Time-triggered:


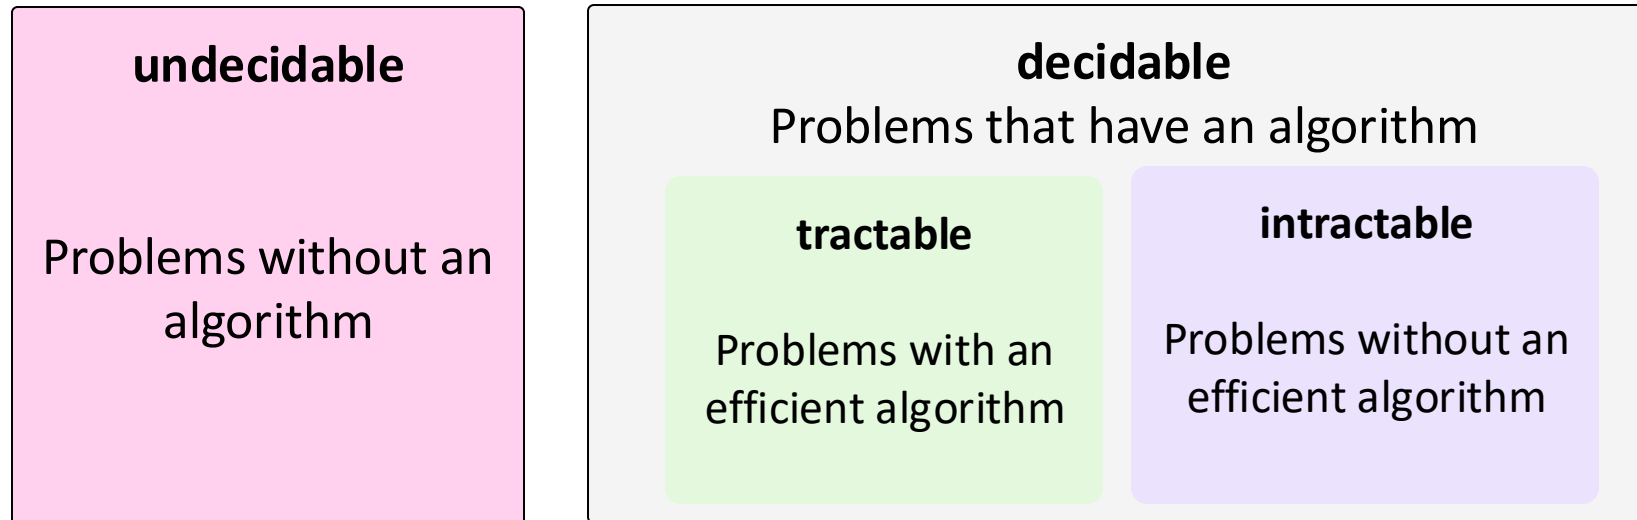
- execution of all other tasks remains the same
- no state change



BUT: Finding the static schedule table is still a very complex problem to solve

# Decidable vs. Computable


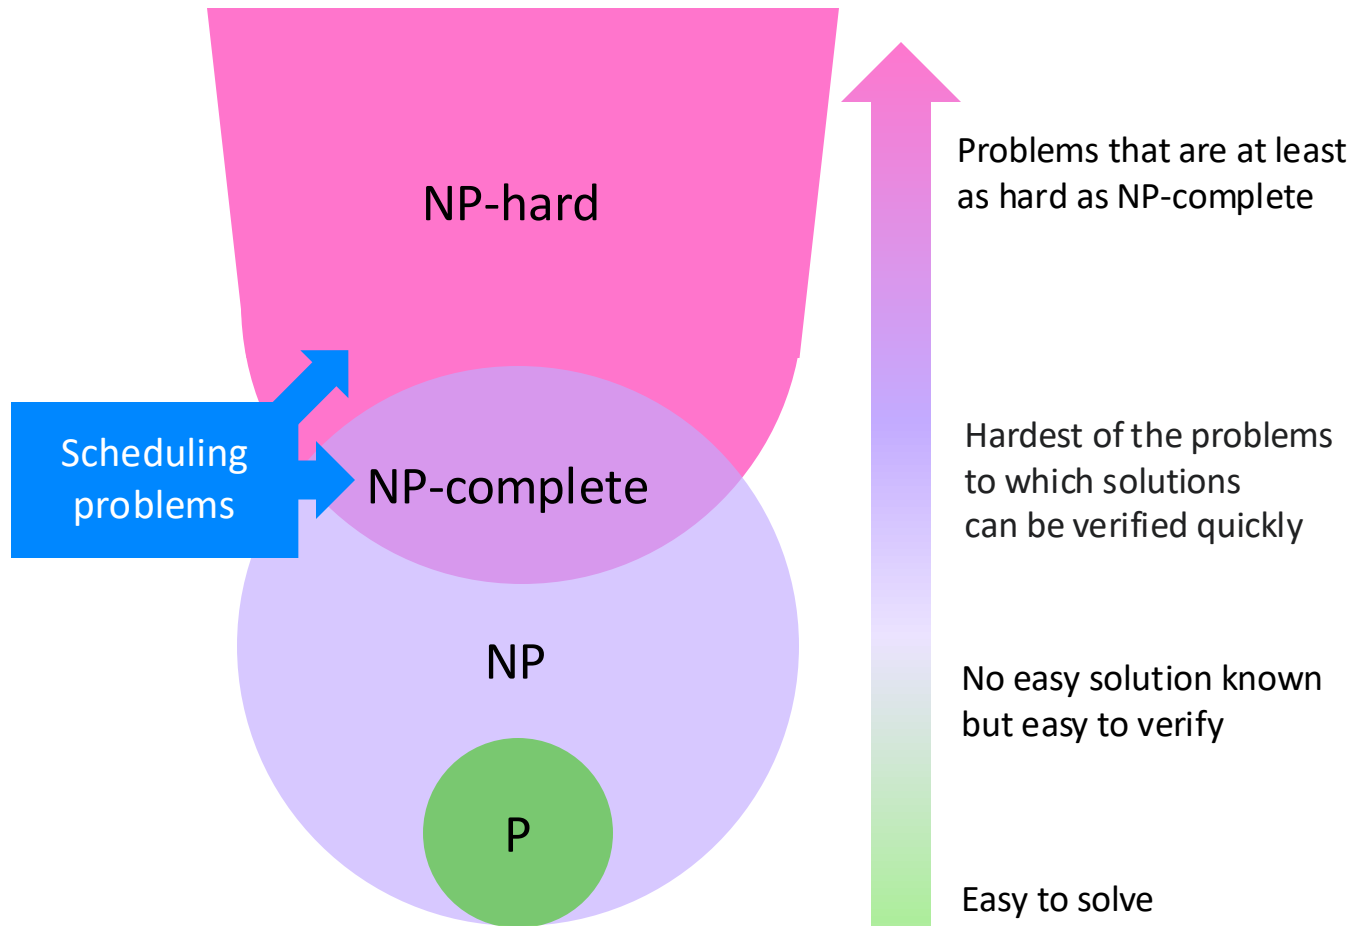
- » **Fundamental question:** David Hilbert asked: is mathematics decidable, i.e., is there a way (algorithm) to prove a statement given a set of axioms and rules?
- » **Alan Turing** answered the question and, in the process, invented the foundation of all modern computers: the universal Turing machine



<https://www.youtube.com/watch?v=HeQX2HjkcNo>

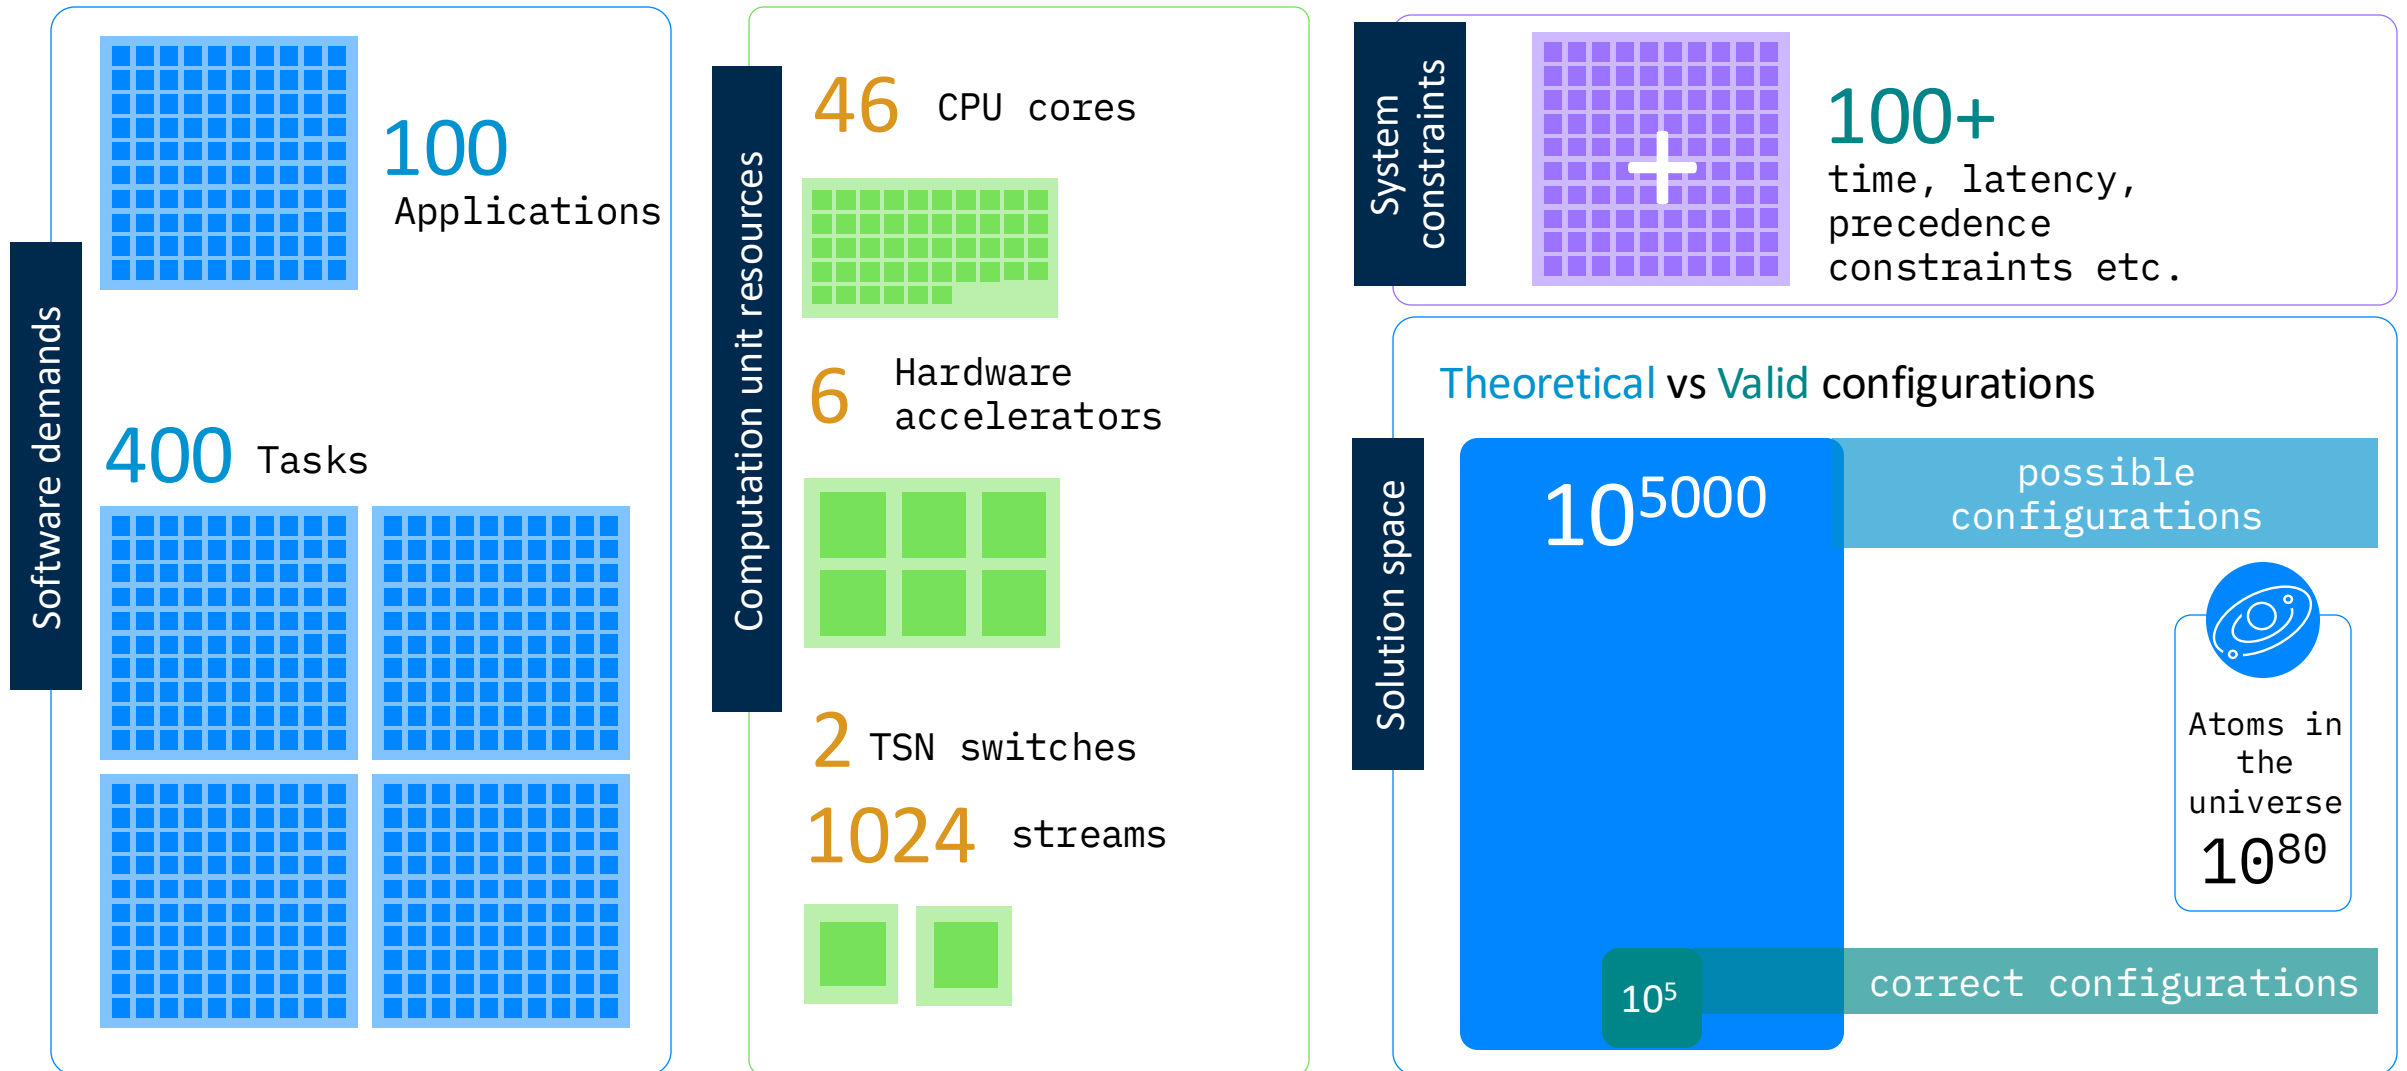
The time-triggered scheduling/configuration problem is decidable.  
**The main question is, is it tractable?**

# Complexity classes




Complexity is not always equal to algorithm efficiency: compare  $1.00001^n$  vs  $n^{50}$

# Complexity in the real-world – ADAS L2 example



# Optimality



## Optimal algorithm:

- » Looks at entire configuration space
- » Finds a solution if there is one (might take a long time)
- » If no solution is found, no other algorithm can find a solution
- » Typically, NP-complete

## Heuristic algorithm:

- » Looks at a fraction of the configuration space (search space)
- » Trade-off performance vs optimality
- » Not guaranteed to find a solution
- » Designed to be fast
- » Tries to find near-optimal solution






# Time-triggered and event-triggered integration

# Downsides of Time-triggered

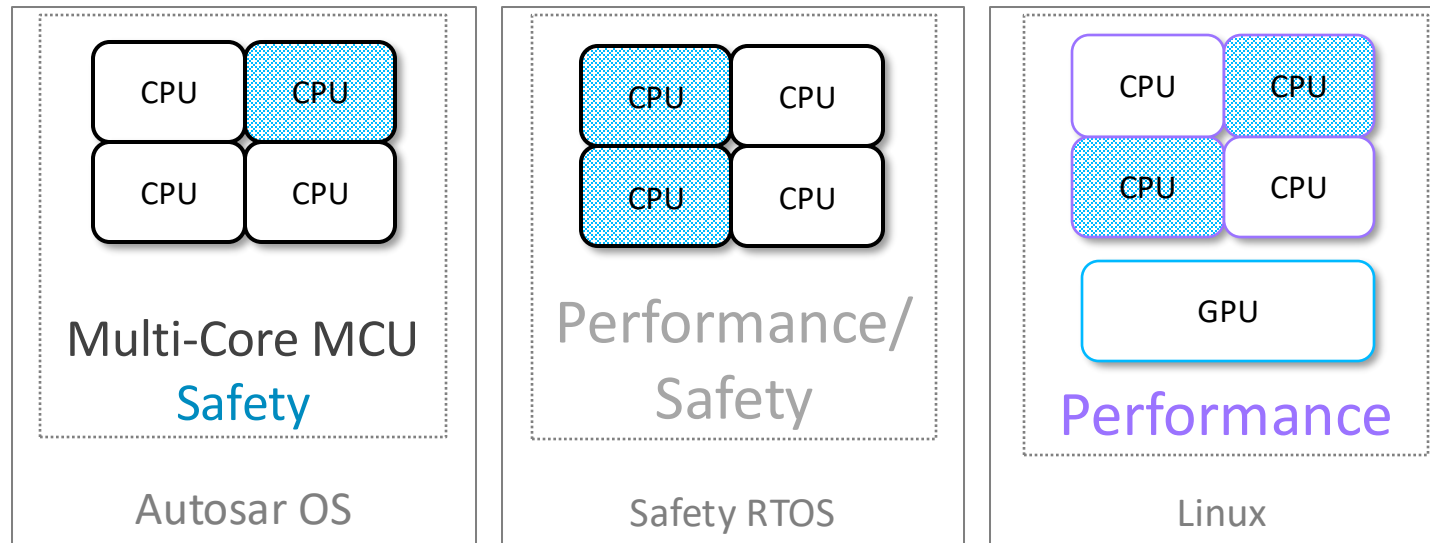
- **Inflexible:** no (big) adaptations at runtime, task execution is fixed
  - event-triggered tasks (sporadic activation)
- **Overprovisioning:** Slots must be reserved for the worst-case
  - WCET analysis is hard and very pessimistic!
- **Precomputed:** task properties must be known at design time
  - this is a development process issue!
- **System limitations:** macrotick, hyperperiod
- **Event Driven tasks:**
  - sporadic activation
  - WCET known
  - constrained or arbitrary deadlines



As much flexibility as possible, as much time-triggered as necessary

# Modern safety-critical systems (e.g., ADAS)

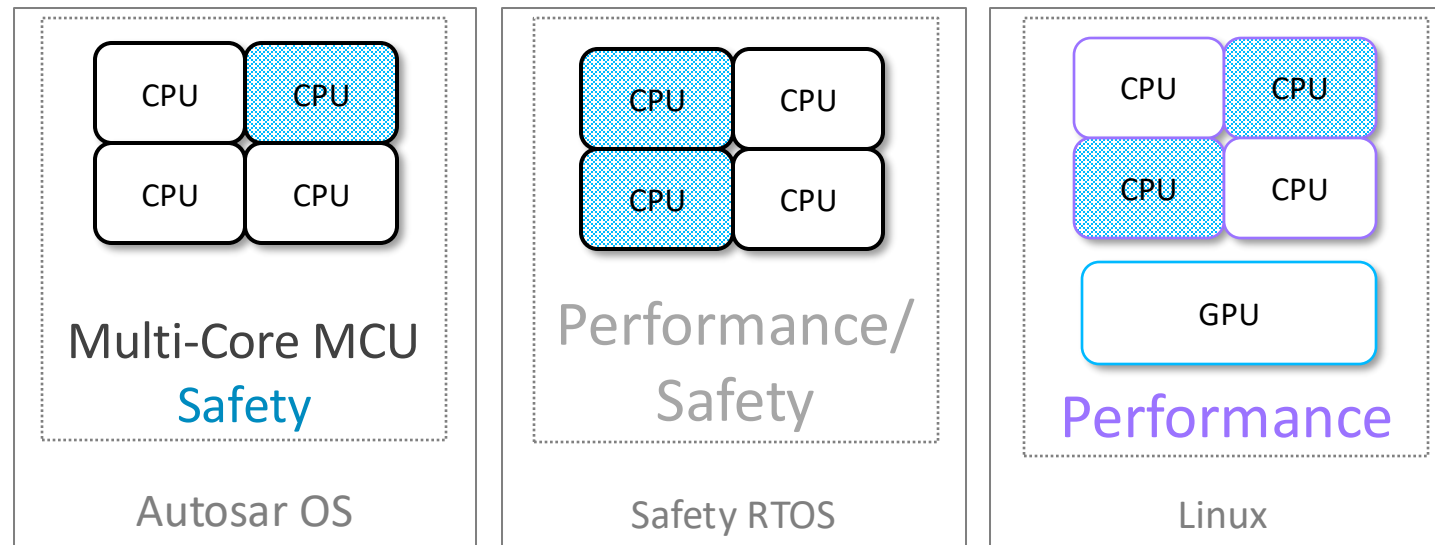
- **Heterogeneous** multi-core multi-SoC platforms featuring a variety of CPUs and GPUs
- Time-Triggered (TT) tasks are **periodic** and **statically scheduled**
- Event-triggered (ET) are **sporadic** but also require **timing guarantees**
- **Partitioned** scheduling approach without runtime migration of tasks



Modern safety-critical systems benefit most from combining TT and ET, allowing a system to be flexible enough to respond to sporadic events when needed, i.e., **best of both worlds**

# Dedicated core


- Dedicate cores to process event-based tasks
- Can be done either using OS scheduler (usually fixed-priority) or via a custom scheduler if needed
- Can also be almost dedicated core : have a few TT slots if needed
- Best for long (non-preemptable) chains that are data-driven and triggered by events



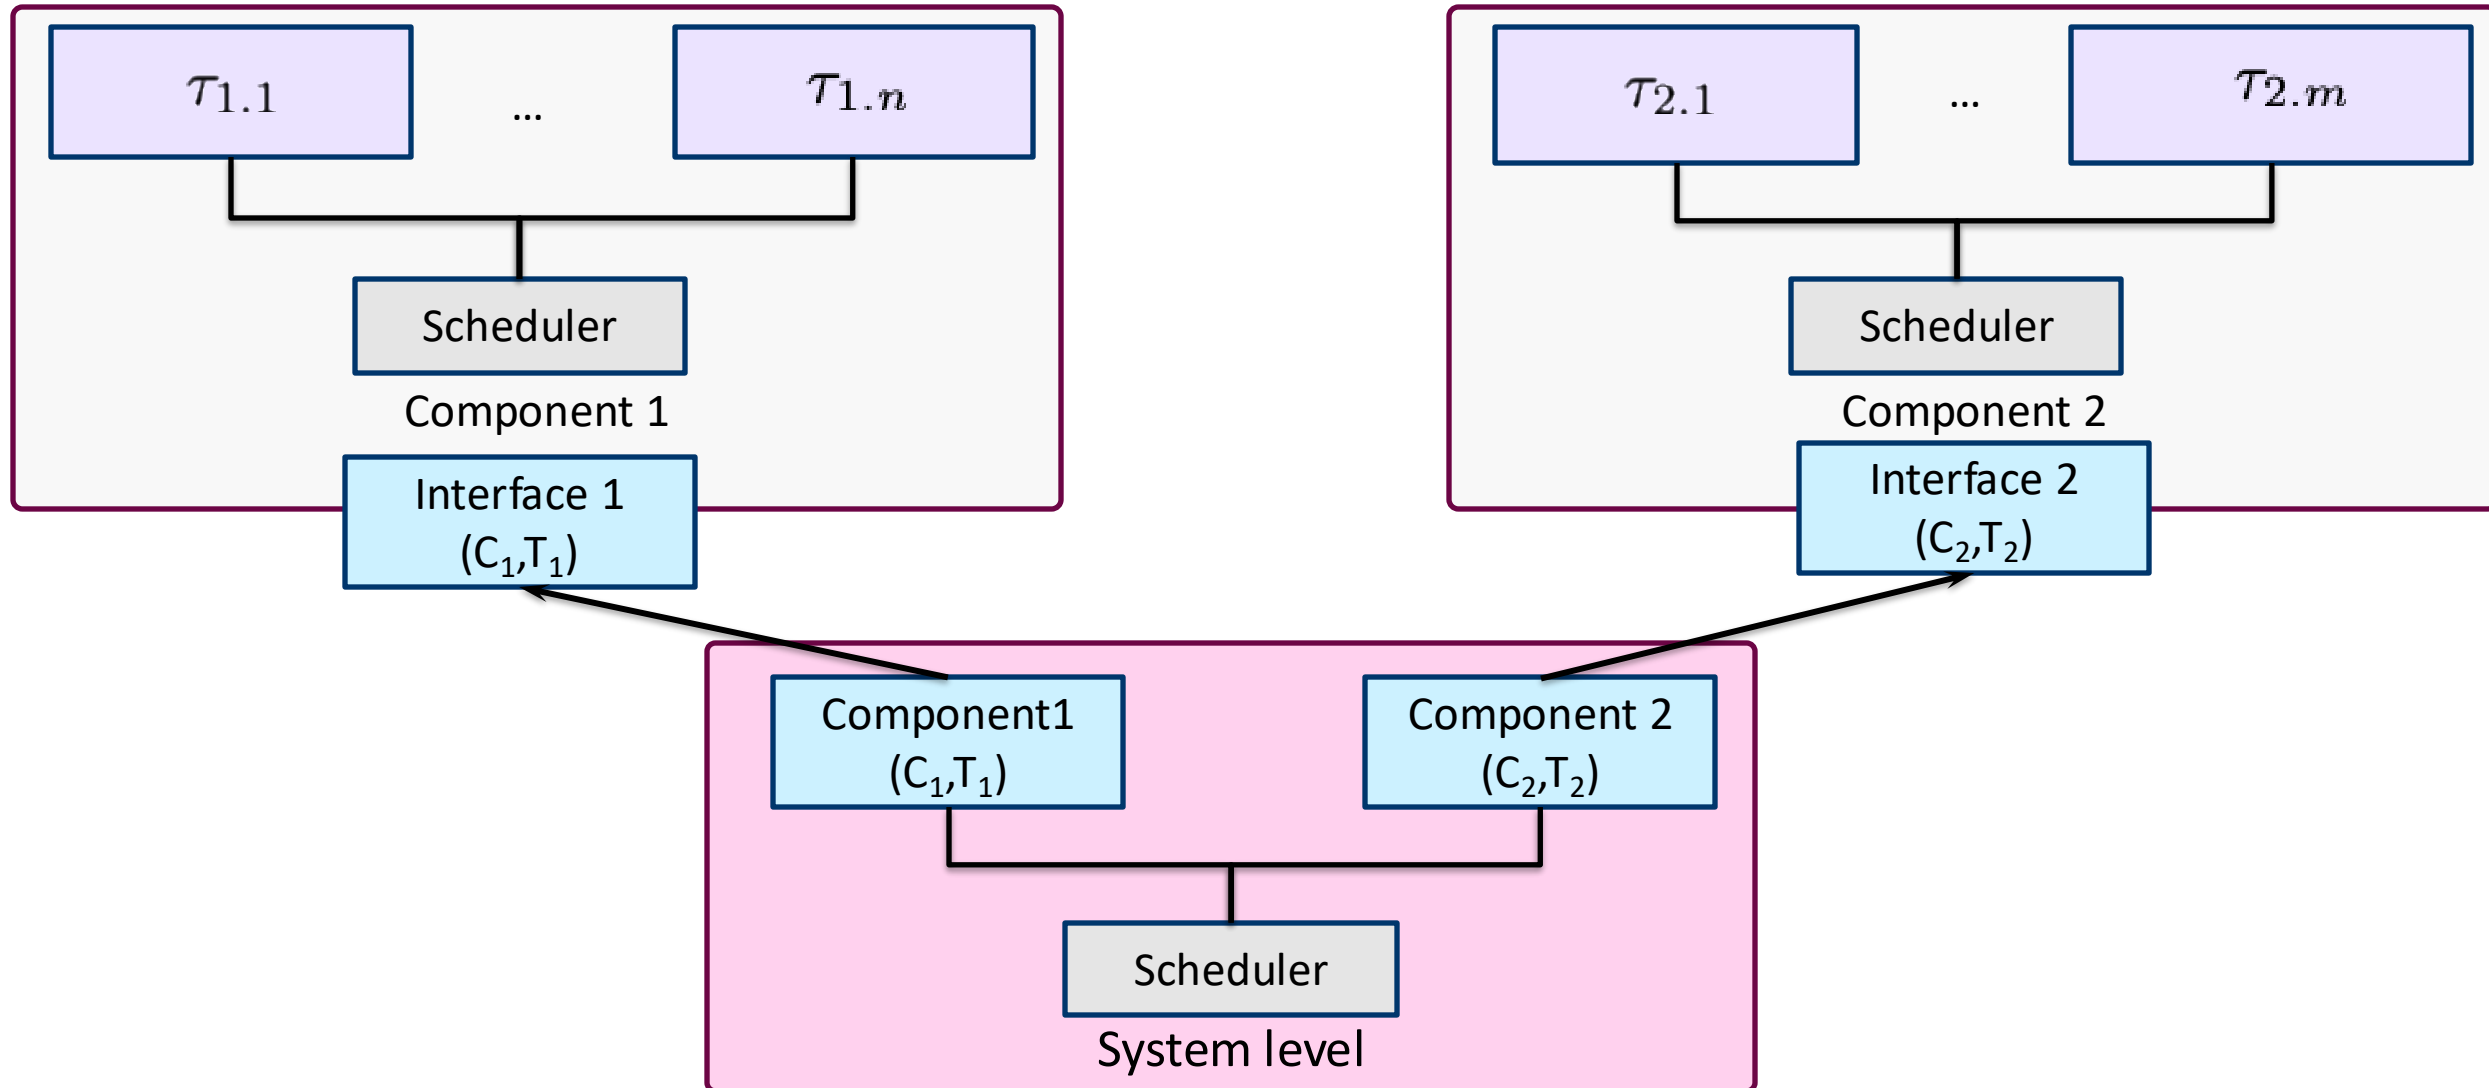
# Slack stealing

**Idea:** Use available slack to process event-based tasks as soon as possible


- Keep track of the system slack (i.e. when tasks finish early), and detect when & if this slack can be used
- Does not modify the TT schedule
- No real guarantees can be given for events
- Ideal for: best-effort event task handling



# Hierarchical scheduling framework



# TT and ET integration

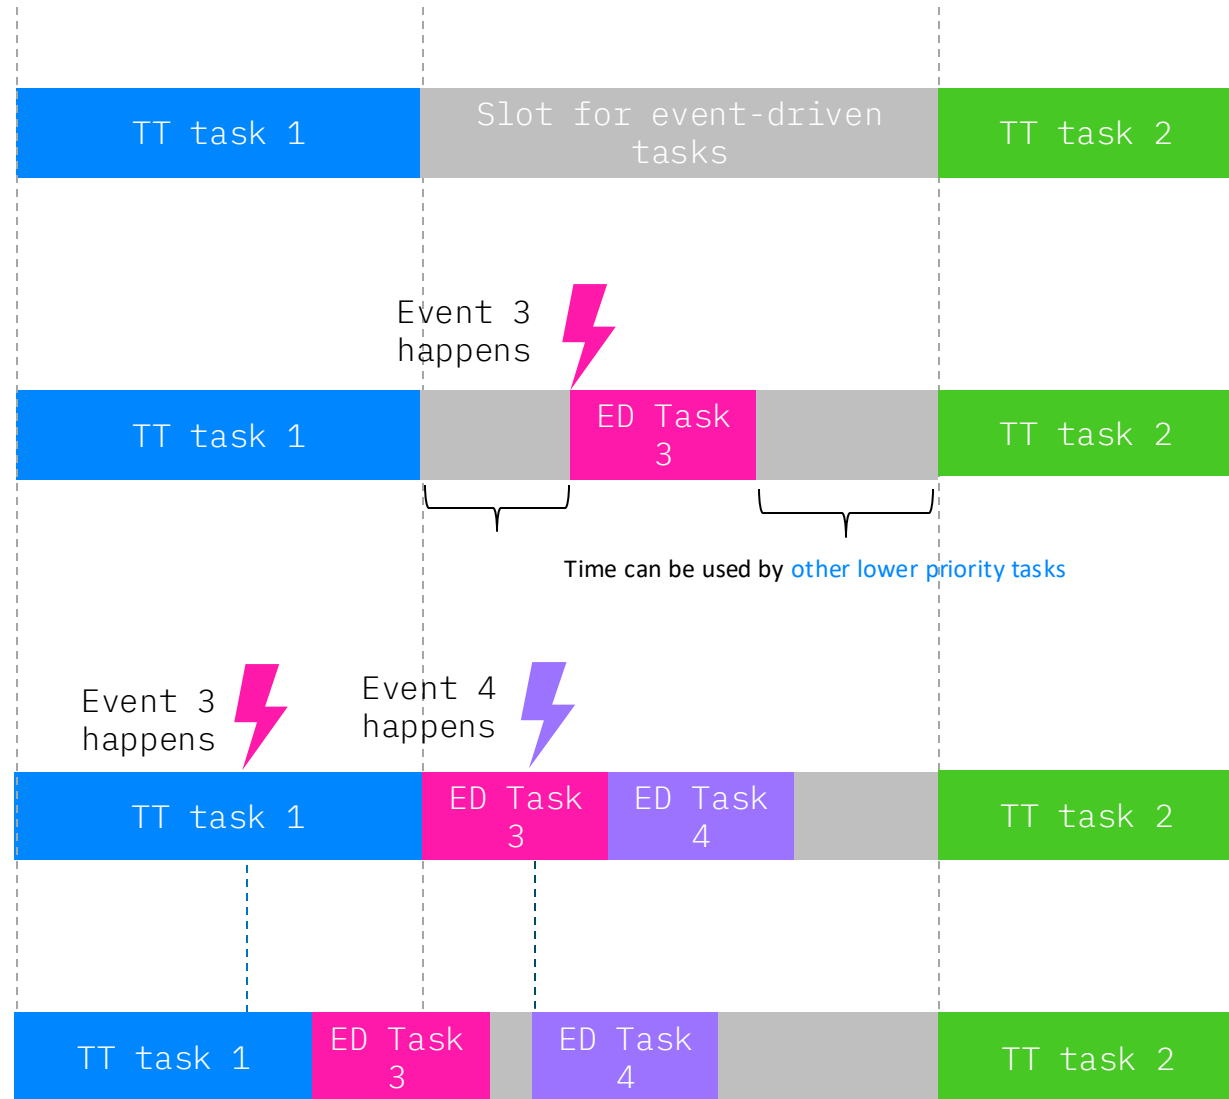


- Time-triggered schedule ensures timing behavior of TT tasks and is calculated offline
- Second-level fixed-priority scheduler that services event-triggered tasks in the idle slots at runtime

# Example Use Cases

Generated schedule table

An event arrives during the slot where event-driven tasks can be executed

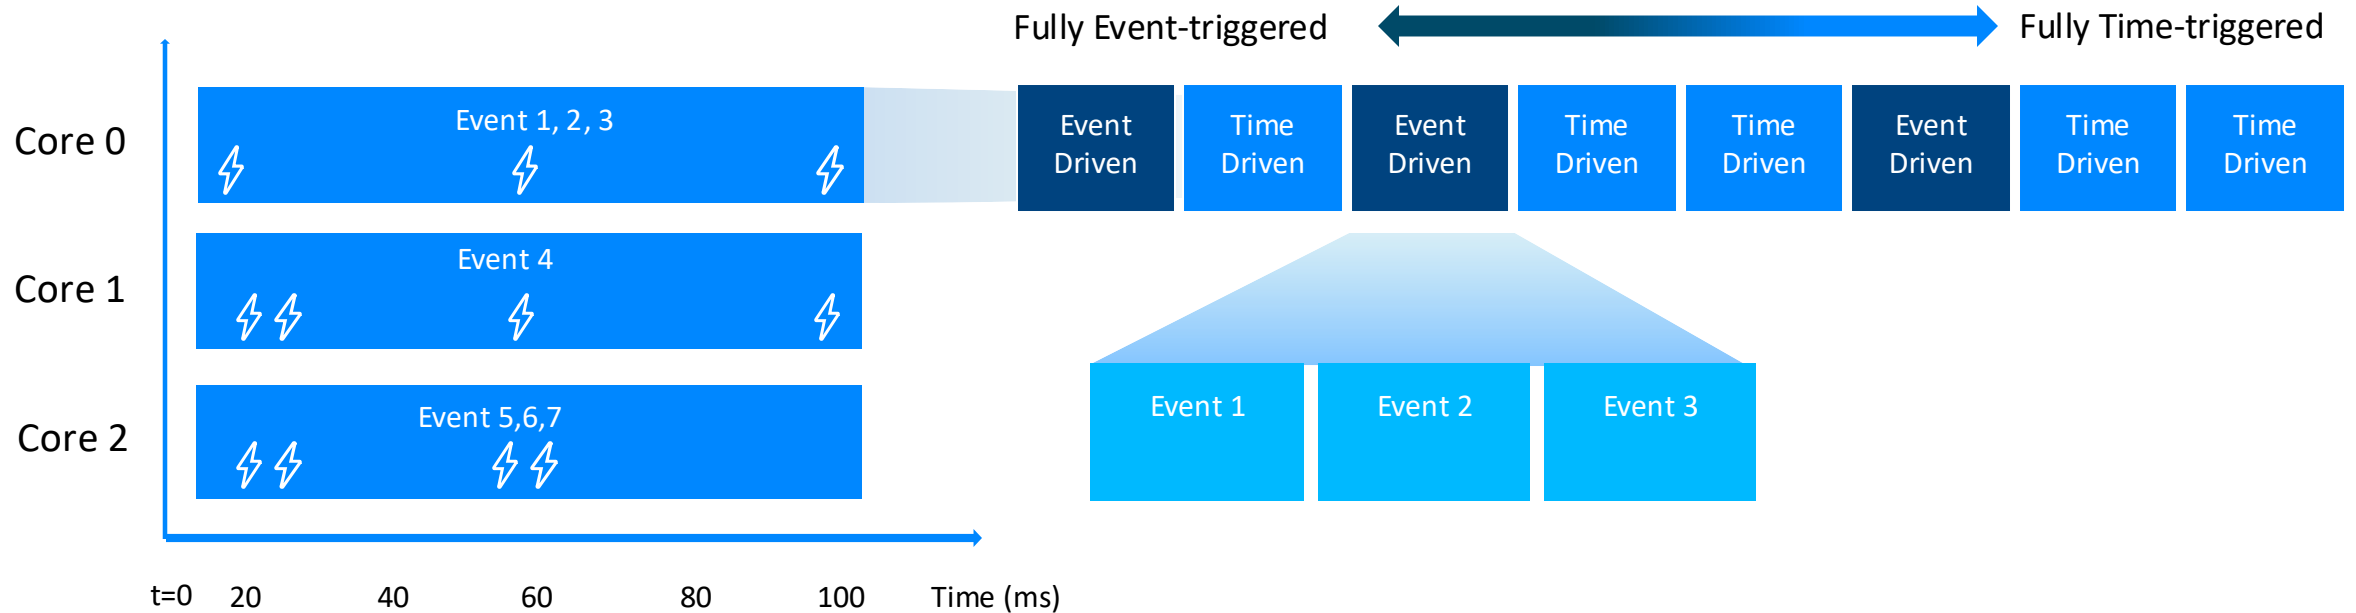


An event arrives during execution of a TT task

The event-driven task **will start in the TT slot** if the TT task finishes earlier !



# TT and ET task allocation and interaction



- The system is partitioned, i.e., TT and ET tasks need to be allocated to cores at design-time
- The TT and ET task allocation problem is crucial to increase schedulability of the system
- The problem is NP-complete, i.e., optimal algorithms are infeasible in the general case

**Problem dimension 1: How to place idle slots such that both TT and ET tasks respect their deadlines?**

**Problem dimension 2: Find a task to core allocation such that problem dimension 1 has a solution (optimization).**

# How can we check ET task schedulability?


Simulate Fixed-Priority and see if deadlines are met

- We need to do this for every potential correct TT schedule candidate (there are a lot of them)
- ...and for every priority ordering (remember the combinatorial explosion of  $n!$ )
- ...and for every arrival pattern of events (also leads to combinatorial explosion)



Arrival of ET



TT schedule



Idle slots




Highly infeasible!


# Oversampling – SPoll

Simple polling: Use one periodic polling task per event – oversampling  
How long is the sampling period?

a) Strictly periodic



b) Non-strictly-periodic



Can be quite inefficient!

Think of an event with  $C = 2$  ms,  $D = 20$  ms, and  $T = 100$ ms. (2% utilization)  
We have to reserve a slot every 9 ms, consuming 22,2 % of CPU bandwidth.

# Solution using response-time analysis

Holistic scheduling [\[Pop2003\]](#)

1. Generate TT schedule such that TT tasks are fulfilled.
2. Check schedulability over resulting idle slots using the response-time method over partitioned resources for every event-based task.
3. Recompute TT schedule if ET tasks not feasible

Slot-shifting [\[Isovic2009\]](#) is a similar method with a different schedulability test that does not assume FP scheduling.

Can be very time-consuming!  
No guidance on how to place idle slots.



Can we do better?

# Periodic resource abstraction

Dedicated resource:

- available all the time at full capacity




Fractional resource:

- available all the time at reduced capacity




Partitioned resource:

- available some of the time at full capacity



Periodic resource:


- available periodically at full capacity (e.g.,  $R(2,3,3)$ )




# Periodic resource abstraction

Explicit Deadline Periodic:  $R(C,D,T)$  [\[Shin2008\]](#)

- We can do the schedulability analysis for ET independently of the TT schedule
- We then treat the polling tasks (periodic resources) as normal TT tasks and create the TT schedule for all
- We still have to find a feasible TT schedule for both TT and polling tasks



# Response-time analysis



$$demand(\tau_i, t) = C_i + \sum_{\forall \tau_j \in HP(\tau_i)} \left\lceil \frac{t}{T_j} \right\rceil \cdot C_j.$$

# Response-time analysis

To verify the schedulability of constrained-deadline task  $\tau_i$  under fixed-priority scheduling in uniprocessor systems, the time-demand analysis (TDA) [Lehoczky89] can be adopted.

Task  $\tau_i$  is schedulable under DM scheduling, where  $HP(\tau_i)$  is the set of tasks with higher priority than task  $\tau_i$  since, if and only if

$$\exists t \text{ with } 0 < t \leq D_i \text{ and } C_i + \sum_{\forall \tau_j \in HP(\tau_i)} \left\lceil \frac{t}{T_j} \right\rceil C_j \leq t$$

Computation time of  $\tau_i$

The number of times task  $\tau_j$  is executed within an interval of length  $t$  in the worst case


Computation time of  $\tau_j$



# Response-time analysis on a dedicated resource

Dedicated resource:

- available all the time at full capacity





$$RT(\tau_i, R) = \text{earliest } t : supply_R(t) \geq demand(\tau_i, t)$$

# Response-time analysis on a partitioned resource

Partitioned resource:

- available some of the time at full capacity




$$RT(\tau_i, R) = \text{earliest } t : supply_R(t) \geq demand(\tau_i, t)$$

# Periodic resource abstraction - AdvPoll

Explicit Deadline Periodic:  $R(C,D,T)$  [Shin2008]

- we do not care when the supply  $C$  is given, as long as it is given in every period  $T$  until the deadline  $D$
- we can use the worst-case linear supply bound and schedulability test defined below [Almeida2004]



$$\alpha = \frac{C}{T}$$

$$\Delta = T + D - 2C$$

$$lslb(t) = \max\{0, (t - \Delta) \cdot \alpha\}$$

$$RT(\tau_i, R) = \text{earliest } t :$$

$$t \geq \Delta + \text{demand}(\tau_i, t) / \alpha$$

$$RT(\tau_i, R) = \text{earliest } t : t \geq \Delta +$$

$$\frac{1}{\alpha} \left( C_i + \sum_{\forall \tau_j \in HP(\tau_i)} \left\lceil \frac{t}{T_j} \right\rceil \cdot C_j \right)$$


# Periodic resource abstraction - AdvPoll

Explicit Deadline Periodic:  $R(C,D,T)$  [Shin2008]

- we do not care when the supply  $C$  is given, as long as it is given in every period  $T$  until the deadline  $D$
- we can use the worst-case linear supply bound and schedulability test defined below [Almeida2004]

- Schedulability analysis for ET is independent of the TT schedule
- The polling tasks (periodic resources) can be scheduled as normal TT tasks

Tradeoff – schedulability for runtime



$$t \geq \Delta + \text{demand}(\tau_i, t) / \alpha$$

$$RT(\tau_i, R) = \text{earliest } t : t \geq \Delta + \frac{1}{\alpha} \left( C_i + \sum_{\forall \tau_j \in HP(\tau_i)} \left\lceil \frac{t}{T_j} \right\rceil \cdot C_j \right)$$

## Another way to express this schedulability condition

Task  $\tau_i$  is schedulable under DM scheduling, where  $HP(\tau_i)$  is the set of tasks with higher priority than task  $\tau_i$  since, if and only if

$$\bigwedge_{i=1 \dots n} \bigvee_{t \in \mathcal{P}_{i-1}(T_i)} \sum_{j=1}^i \left\lfloor \frac{t}{T_j} \right\rfloor C_j \leq t$$

$$\mathcal{P}_i(t) \rightarrow \begin{cases} \mathcal{P}_0(t) = \{t\} \\ \mathcal{P}_i(t) = \mathcal{P}_{i-1} \left( \left\lfloor \frac{t}{T_i} \right\rfloor T_i \right) \cup \mathcal{P}_{i-1}(t). \end{cases}$$

Has to hold for all  $\tau_i$

And at least for one of the values in the set of points of interest

## Another way to express this schedulability condition

Earliest point of  $P_{i-1}(D_i)$  for which the workload at  $D_i$  is smaller than  $D_i$

$$C_i + \min_{t \in P_{i-1}(D_i)} \sum_{j=1}^{i-1} \left\lceil \frac{t}{T_j} \right\rceil C_j + (D_i - t) \leq D_i.$$

I am getting only a fraction of the CPU, so scale all execution times

$$\bigwedge_{i=1 \dots n} \bigvee_{t \in P_{i-1}(T_i)} \sum_{j=1}^i \left\lceil \frac{t}{T_j} \right\rceil \frac{C_j}{\alpha} \leq t$$

The tasks are not receiving any execution time from the global scheduler during the interval  $\Delta > 0$

$$\Delta + \frac{C_i}{\alpha} + \min_{t \in P_{i-1}(D_i)} \sum_{j=1}^{i-1} \left\lceil \frac{t}{T_j} \right\rceil \frac{C_j}{\alpha} + (D_i - t) \leq D_i$$

The schedulability condition then becomes:

$$\bigwedge_{i=1 \dots n} \bigvee_{t \in P_{i-1}(D_i)} \Delta \leq t - \frac{1}{\alpha} \sum_{j=1}^i \left\lceil \frac{t}{T_j} \right\rceil C_j.$$

## Another way to express this schedulability condition


$$\bigwedge_{i=1..n} \bigvee_{t \in \mathcal{P}_{i-1}(D_i)} \Delta \leq t - \frac{1}{\alpha} \sum_{j=1}^i \left\lceil \frac{t}{T_j} \right\rceil C_j.$$

Can be put in terms of optimization problem with e.g. switching cost

$$c_1 \frac{T_{\text{Overhead}}}{P} + c_2 \alpha$$

When choosing the server parameters, we must balance two opposite needs:


- the required bandwidth should be small, to not waste the total processor capacity;
- the server period should be large, otherwise the time wasted in context switches performed by the global scheduler will be too high.





# Single-core experiments

- 30 TT and 20 ET tasks per task set
- Periods  $\in \{5, 10, 20, 40, 80\}$  ms
- Microtick of 250 $\mu$ s
- Constrained Deadline:  $D_i$  is uniformly selected in upper half of  $[C_i, T_i]$
- Arbitrary Deadline:  $D_i \in [C_i, 5 \cdot T_i]$
- 20% TT task utilization
- Increasing ET utilization
- 100 task sets per configuration



**SPoll** is quickest and can be used for low utilization

**Holistic** has better schedulability at high utilization but takes much longer time

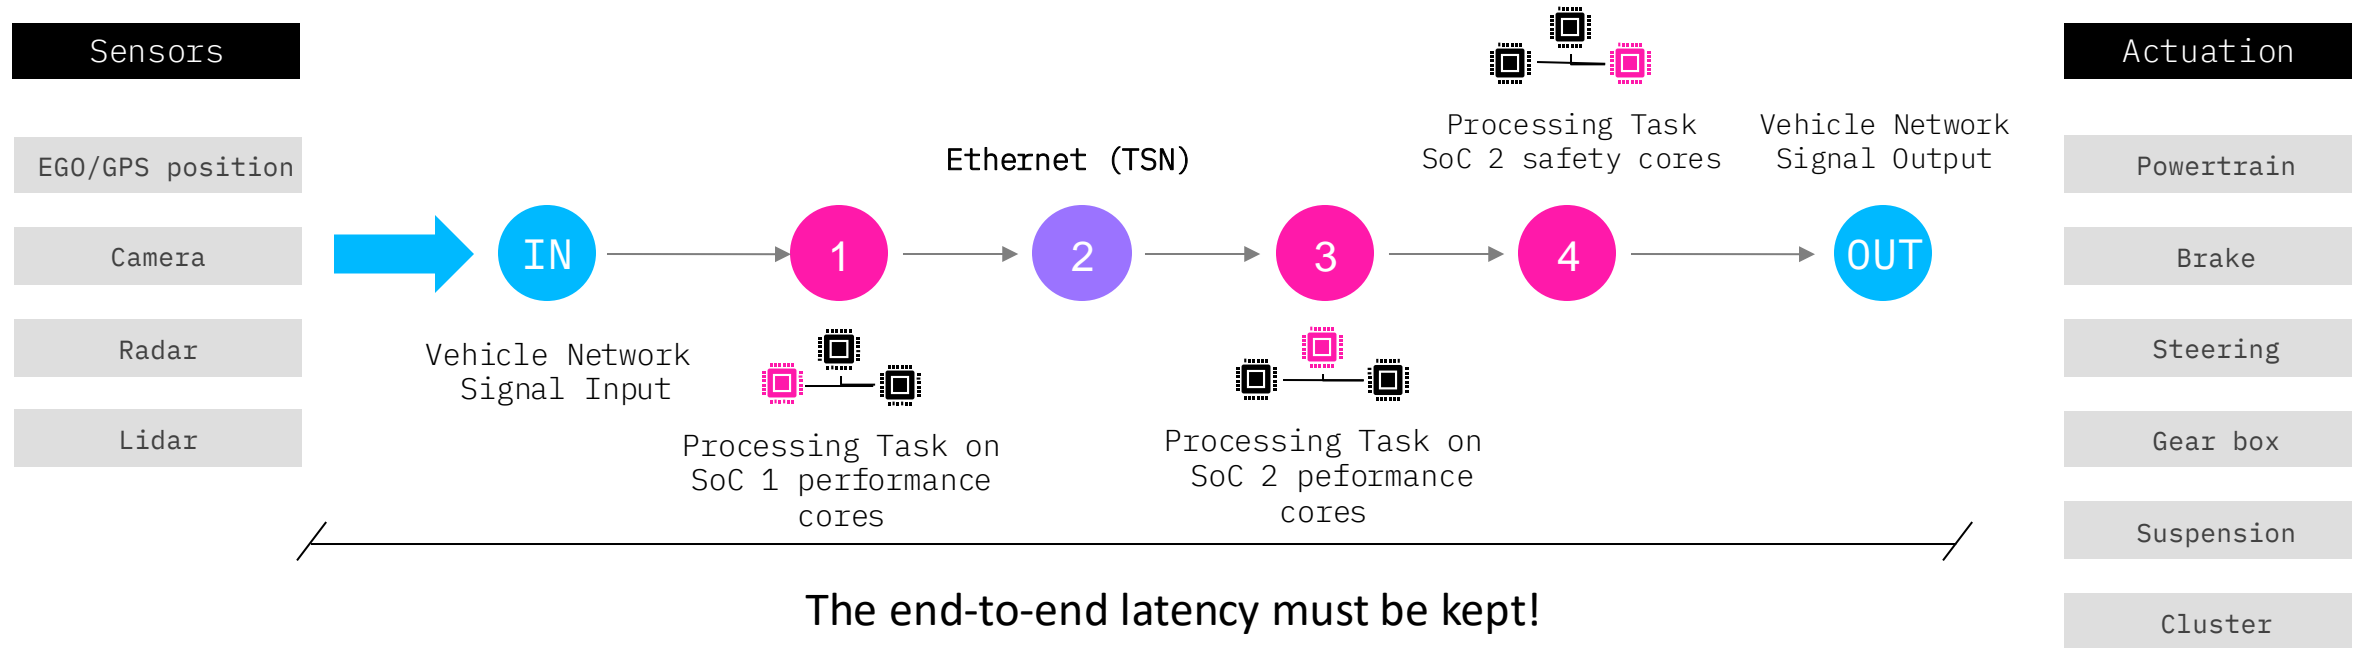
**AdvPoll** has the best overall trade-off between runtime and schedulability



Real-time networks


# Computation Chains Spanning across a Network of Chips

Example: ADAS/AD workload




> Task Scheduling and Network Communication must be aligned across SoCs and CPUs

# Automotive networks




**n ECU: n functions**

DISTRIBUTED  
E/E ARCHITECTURE



**5-7 domains: n functions**

DOMAIN  
E/E ARCHITECTURE



**2 ECUs: n functions**

CENTRALISED  
E/E ARCHITECTURE

YESTERDAY & TODAY

2020 – 2025

2025 +

# Automotive networks


## » Where do we come from?

- Increased number of ECUs led to lots of wiring.

DAIMLER

## Mercedes-Benz S-Class (2006) complete cable harness

- about 38kg, following slides show details of this harness (parts of some dedicated harnesses)

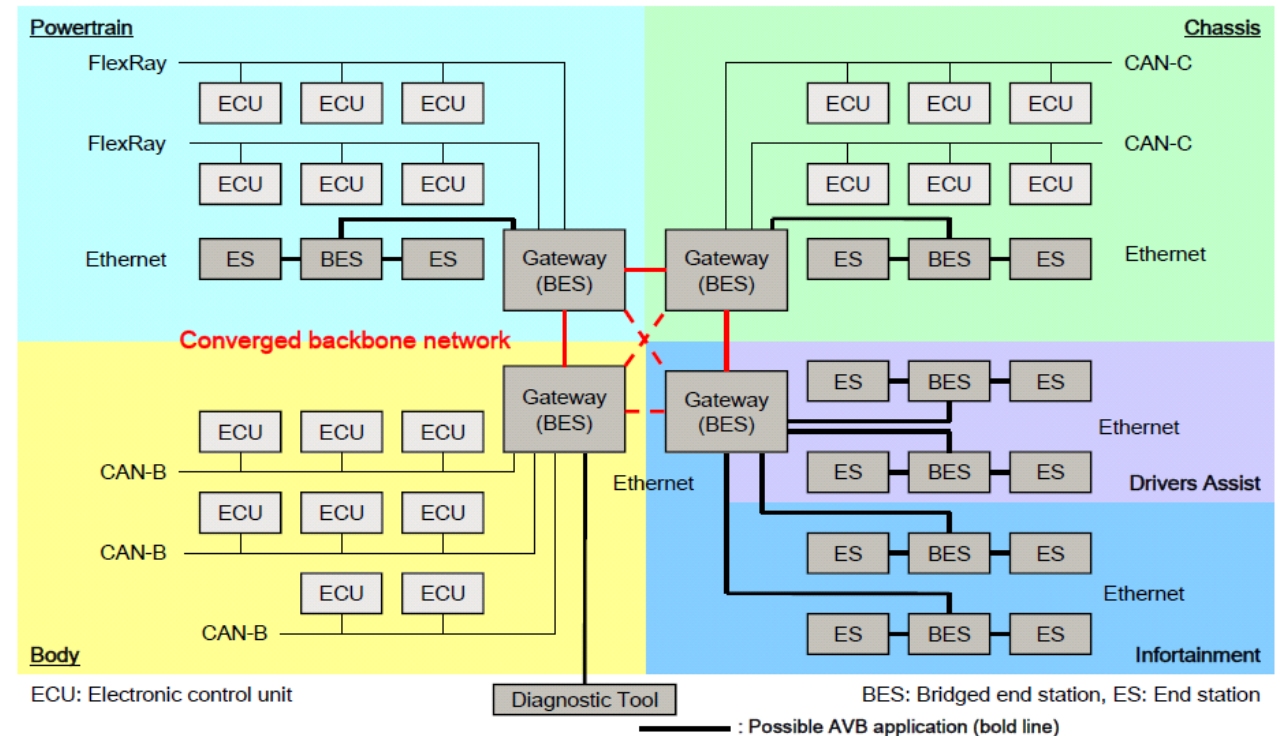


Source: IEEE 802.3 RTPGE SG

# Automotive networks

## » Where do we come from?

- Increased number of ECUs led to lots of wiring.
- Automotive networks were a mixture of domain and zonal architectures interconnected by different protocols and technologies e.g., CAN, FlexRay and




Source: IEEE 802.1 AVB TG presentation



# Automotive networks

## » Where do we come from?

- Increased number of ECUs led to lots of wiring.
- Communication networks were a mixture of zonal, domain with network technologies with several technologies implied. E.g., CAN, FlexRay and LIN.
- Centralized computing platforms like ICCP aim at gathering the formerly distributed computations into a more powerful unit.
  - Centralization might help to harmonize the communication network.
  - But applications are becoming more complex and might have additional communication requirements.




# Distributed real-time system

Real-time distributed systems

- Guarantee on timely and predictable delivery of messages across the network
- Task execution (message production and consumption) synchronized with communication


## EVENT-DRIVEN SCHEDULING

Event-handlers with predictable response time




## TIME-TRIGGERED SCHEDULING

For resource allocation / partitioning, typically done by Software / System Integrator



## NETWORK SCHEDULING

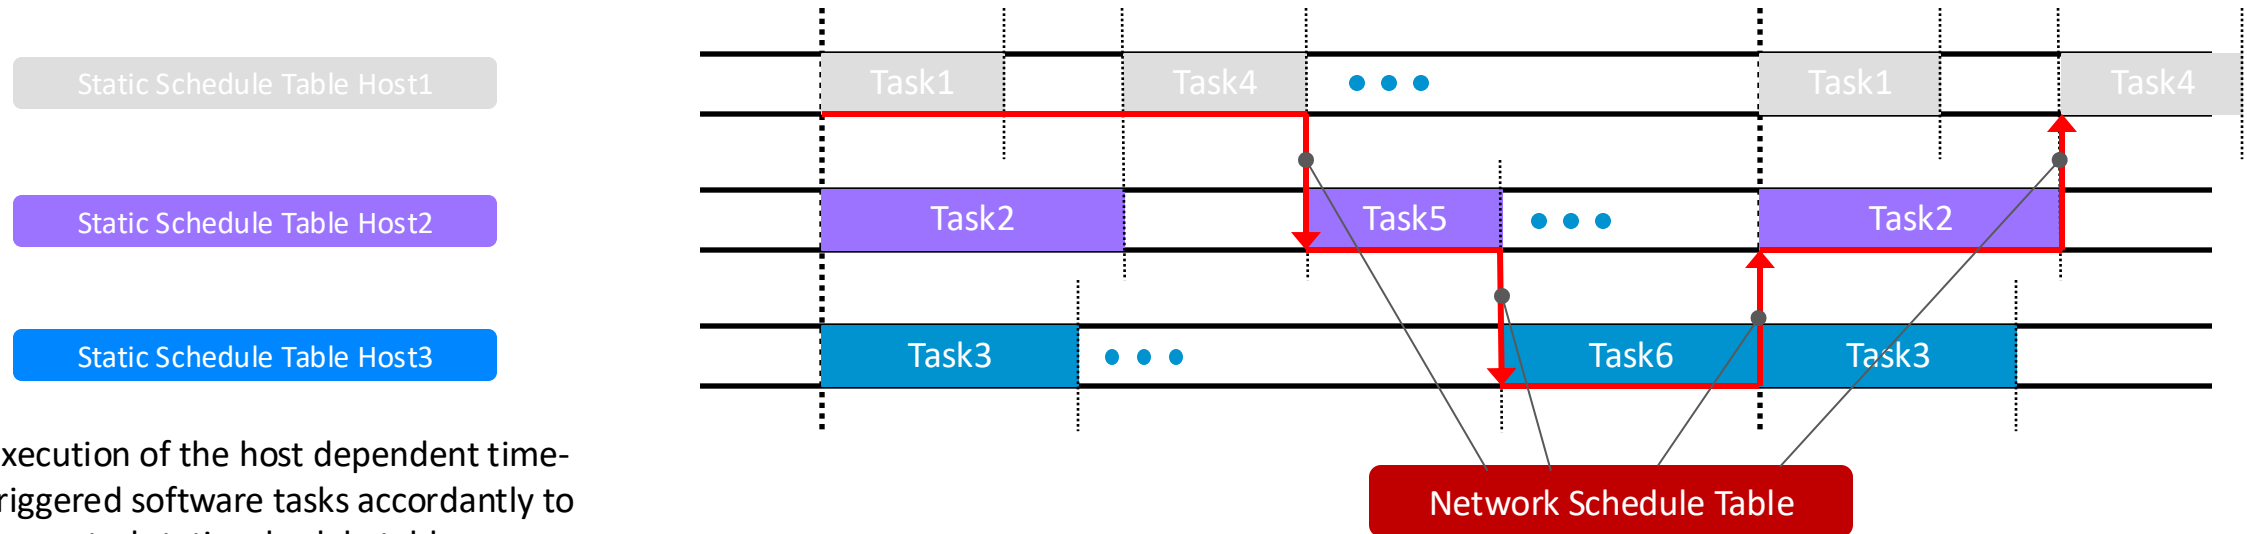
The network scheduling, enables to ensure end-to-end behaviors across SoCs





# Time-Triggered scheduling


## Global Schedule Algorithm



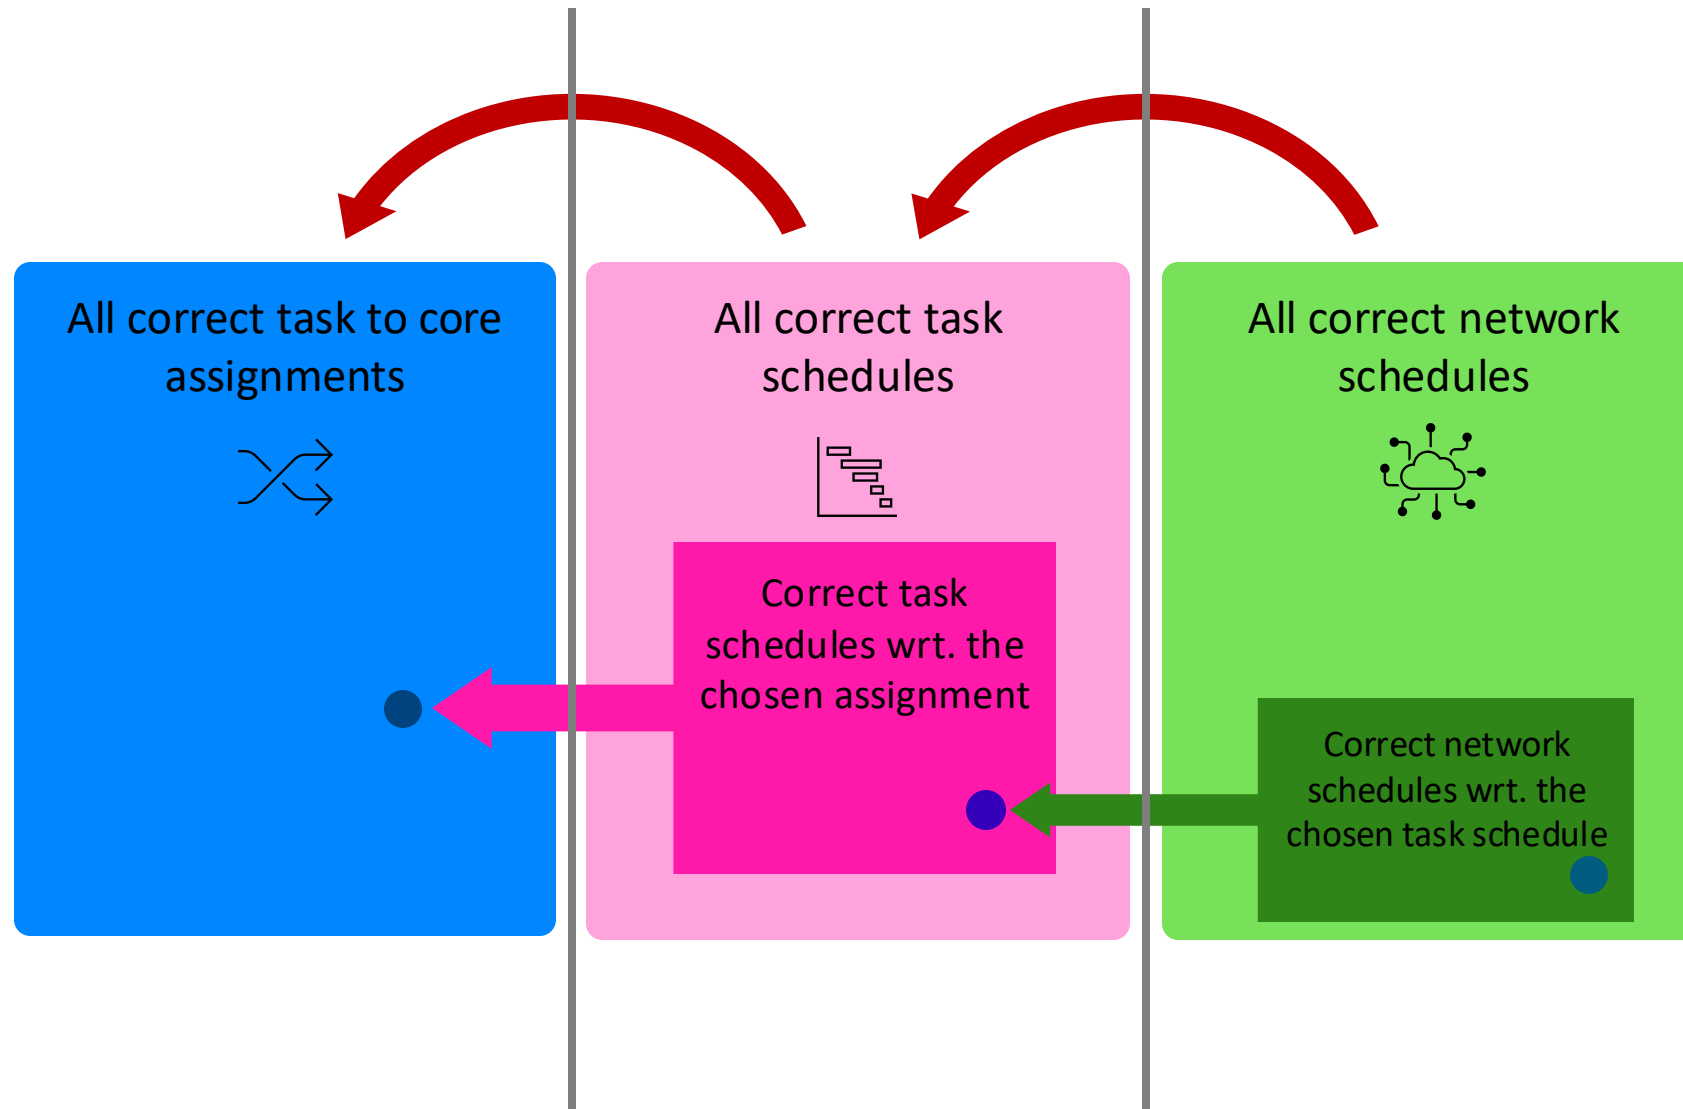
Execution of the host dependent time-triggered software tasks accordantly to generated static schedule tables

Handling of the communication messages accordantly to generated time-triggered timing

# Problem dimensions



# Problem dimensions



# Thank you!



[silviu.craciunas@tttech.com](mailto:silviu.craciunas@tttech.com)



<https://scraciunas.github.io/>



[www.linkedin.com/in/scraciunas](https://www.linkedin.com/in/scraciunas)